

Deterministic and stochastic inexact regularization algorithms for nonconvex optimization with optimal complexity

S. Bellavia*, G. Gurioli†, B. Morini‡ and Ph. L. Toint§

7 November 2018

Abstract

A regularization algorithm using inexact function values and inexact derivatives is proposed and its evaluation complexity analyzed. This algorithm is applicable to unconstrained problems and to problems with inexpensive constraints (that is constraints whose evaluation and enforcement has negligible cost) under the assumption that the derivative of highest degree is β -Hölder continuous. It features a very flexible adaptive mechanism for determining the inexactness which is allowed, at each iteration, when computing objective function values and derivatives. The complexity analysis covers arbitrary optimality order and arbitrary degree of available approximate derivatives. It extends results of Cartis, Gould and Toint (2018) on the evaluation complexity to the inexact case: if a q th order minimizer is sought using approximations to the first p derivatives, it is proved that a suitable approximate minimizer within ϵ is computed by the proposed algorithm in at most $O(\epsilon^{-\frac{p+\beta}{p-q+\beta}})$ iterations and at most $O(|\log(\epsilon)|\epsilon^{-\frac{p+\beta}{p-q+\beta}})$ approximate evaluations. While the proposed framework remains so far conceptual for high degrees and orders, it is shown to yield simple and computationally realistic inexact methods when specialized to the unconstrained and bound-constrained first- and second-order cases. The deterministic complexity results are finally extended to the stochastic context, yielding adaptive sample-size rules for subsampling methods typical of machine learning.

Keywords: evaluation complexity, regularization methods, inexact functions and derivatives, subsampling methods, machine learning.

1 Introduction

Evaluation complexity of algorithms for nonlinear and possibly nonconvex optimization problems has been the subject of active research in recent years. This field is concerned by deriving formal bounds on the number of evaluations of the objective function (and possibly of its derivatives) necessary to obtain approximate optimal solutions within a user-specified

*Department of Industrial Engineering of Florence, Università degli Studi, Firenze, Italy. Member of the INdAM Research Group GNCS. Email: stefania.bellavia@unifi.it

†Department of Industrial Engineering of Florence, Università degli Studi, Firenze, Italy. Email: gianmarco.gurioli@unifi.it

‡Department of Industrial Engineering of Florence, Università degli Studi, Firenze, Italy. Member of the INdAM Research Group GNCS. Email: benedetta.morini@unifi.it

§Namur Center for Complex Systems (naXys), University of Namur, 61, rue de Bruxelles, B-5000 Namur, Belgium. Email: philippe.toint@unamur.be

accuracy. Until recently, the results had focused on methods using first- and second-order derivatives of the objective function, and on convergence guarantees to first- or second-order stationary points [24, 19, 20, 15, 8]. Among these contributions, [20, 8] analyzed the “regularization method”, in which a model of the objective function around a given iterate is constructed by adding a regularization term to the local Taylor expansion, model which is then approximately minimized in an attempt to find a new point with a significantly lower objective function value [17]. Such methods have been shown to possess optimal evaluation complexity [10] for first- and second-order models and minimizers, and have generated considerable interest in the research community. A theoretically significant step was made in [4] for unconstrained problems, where evaluation complexity bounds were obtained for convergence to first-order stationary points of a simplified regularization methods using models of arbitrary degree. Even more recently, [9] proposed a conceptual unified framework subsuming all known results for regularization methods, establishing an upper evaluation complexity bound for arbitrary model degree and also, for the first time, for arbitrary orders of optimality. This paper additionally covers unconstrained problems and problems involving “inexpensive” constraints, that is constraints whose evaluation/enforcement cost is negligible compared to that of evaluating the objective function and its derivatives. It also allows for a full range of smoothness assumptions on the objective function. Finally it proves that the complexity results obtained are optimal in the sense that upper and lower evaluation complexity bounds match in order. In [9], all the above mentioned results are established for versions of the regularization algorithms where it is assumed that objective function values and values of its derivatives (when necessary) can be computed exactly.

In practice, it may sometimes be difficult or impossible to obtain accurate values of the problem’s function and/or derivatives. This difficulty has been known for a long time and has generated its own stream of results, among which we note the trust-region method using dynamic accuracy on the objective function and (possibly on) its gradient (see Sections 8.4.1.1 and 10.6 of [13] and [1]), and the purely probabilistic approaches of [21] and [5]. Since unconstrained cubic regularization methods have become popular in the machine learning community due to their optimal complexity, several contributions have considered building those function and derivative’s approximations by ‘subsampling’ the (very many) nonlinear terms whose sum defines the objective functions typical of machine learning applications. Inexact Hessian information is considered in [12, 2, 25, 26], approximate gradient and Hessian evaluations are used in [11, 22, 27], function, gradient and Hessian values are sampled in [18, 3]. The amount of inexactness allowed is controlled dynamically in [11, 18, 12, 2].

Contributions. The present paper proposes an extension of the unifying framework of [9] for unconstrained or inexpensively-constrained problems that allows inexact evaluations of the objective function and of the required derivatives, in an adaptive way inspired by the trust-region scheme of [13, Section 10.6]. This extension has the advantage of preserving the optimal complexity of the accurate regularization methods and, as in [9], evaluation complexity results are provided for arbitrary model degree and arbitrary order of optimality. In particular, the proposed framework allows all combinations of exact/inexact objective functions and derivatives of any order (including of course degrees and orders one and two, for which simple specializations are outlined). We also consider a stochastic version of our framework and derive rules for sample size in the context of subsampling methods for machine learning.

The paper is organized as follows. Section 2 recalls the notions of high-order optimality proposed in [9] and introduces the general adaptive regularization algorithm allowing dynamic

accuracy (AR p DA). The details of how to obtain the desired relative accuracy levels from known absolute errors are examined in Section 3. The evaluation complexity of obtaining approximate minimizers using this algorithm is then analyzed in Section 4. The general framework is specialized to first- and second-order optimization in Section 5, showing that practical implementation for low order is simple and computationally realistic. The stochastic evaluation complexity and sampling rules for machine learning applications are finally derived in Section 6. Conclusions and perspectives are presented in Section 7.

Notations. Unless otherwise specified, $\|\cdot\|$ denotes the standard Euclidean norm for vectors and matrices. For a general symmetric tensor S of order p , we define

$$\|S\|_{[p]} \stackrel{\text{def}}{=} \max_{\|v\|=1} |S[v]^p| = \max_{\|v_1\|=\dots=\|v_p\|=1} |S[v_1, \dots, v_p]| \quad (1.1)$$

the induced Euclidean norm. We also denote by $\nabla_x^j f(x)$ the j -th order derivative tensor of f evaluated at x and note that such a tensor is always symmetric for any $j \geq 2$. $\nabla_x^0 f(x)$ is a synonym for $f(x)$. $\lceil \alpha \rceil$ and $\lfloor \alpha \rfloor$ denote the smallest integer not smaller than α and the largest integer not exceeding α , respectively. If i is a non-negative integer and β a real in $(0, 1]$ we define $(i + \beta)! = \prod_{\ell=1}^i (\ell + \beta)$. For symmetric matrices, $\lambda_{\min}[M]$ is the leftmost eigenvalue of M . $Pr[\text{event}]$ finally denotes the probability of an event. Finally $\text{globmin}_{x \in \mathcal{S}} f(x)$ denotes the smallest value of $f(x)$ over $x \in \mathcal{S}$.

2 High-order necessary conditions and the AR p DA algorithm

Given $p \geq 1$, we consider the set-constrained optimization problem

$$\min_{x \in \mathcal{F}} f(x), \quad (2.1)$$

where $\mathcal{F} \subseteq \mathbb{R}^n$ is closed and nonempty, and where we assume that the *values of the objective function f and its derivatives must be computed inexactly*. We also assume that $f \in \mathcal{C}^{p,\beta}(\mathbb{R}^n)$, meaning that:

- f is p -times continuously differentiable,
- f is bounded below by f_{low} , and
- the p -th derivative tensor of f at x is globally Hölder continuous, that is, there exist constants $L \geq 0$ and $\beta \in (0, 1]$ such that, for all $x, y \in \mathbb{R}^n$,

$$\|\nabla_x^p f(x) - \nabla_x^p f(y)\|_{[p]} \leq L \|x - y\|^\beta. \quad (2.2)$$

The more standard case where f is assumed to have Lipschitz-continuous p -th derivative is recovered by setting $\beta = 1$ in the above assumptions (for example, the choices $p = 2$ and $\beta = 1$ correspond to the assumption that f has a Lipschitz continuous Hessian). In what follows, we assume that β is known.

If we denote the p th degree Taylor expansion of f around x evaluated at s by

$$T_p^f(x, s) \stackrel{\text{def}}{=} f(x) + \sum_{\ell=1}^p \frac{1}{\ell!} \nabla_x^\ell f(x)[s]^\ell, \quad (2.3)$$

we may then define the *Taylor increment* by

$$\Delta T_p^f(x, s) = T_p^f(x, 0) - T_p^f(x, s). \quad (2.4)$$

Under the above assumptions, we recall the crucial bounds on differences between f and its derivatives and their Taylor's expansion.

Lemma 2.1 [9, Lemma 2.1] Let $f \in C^{p,\beta}(\mathbb{R}^n)$, and $T_p^f(x, s)$ be the Taylor approximation of $f(x + s)$ about x given by (2.3). Then for all $x, s \in \mathbb{R}^n$,

$$f(x + s) \leq T_p^f(x, s) + \frac{L}{(p + \beta)!} \|s\|^{p+\beta}, \quad (2.5)$$

$$\|\nabla_x^j f(x + s) - \nabla_s^j T_p^f(x, s)\|_{[j]} \leq \frac{L}{(p - j + \beta)!} \|s\|^{p+\beta-j}. \quad (j = 1, \dots, p). \quad (2.6)$$

We also follow [9] and define an q th-order-necessary minimizer as a point $x \in \mathbb{R}^n$ such that, for some $\delta \in (0, 1]$,

$$\phi_{f,q}^\delta(x) \stackrel{\text{def}}{=} f(x) - \underset{\substack{x+d \in \mathcal{F} \\ \|d\| \leq \delta}}{\text{globmin}} T_q^f(x, d) = 0. \quad (2.7)$$

Observe that, in the unconstrained case, this definition subsumes the usual optimality criteria for orders one and two, since, if $q = 1$, (2.7) gives that, for any $\delta \in (0, 1]$ (and in particular for $\delta = 1$),

$$\phi_{f,q}^\delta(x) = \|\nabla_x f(x)\| \delta, \quad (2.8)$$

and first-order optimality is thus equivalent to

$$\|\nabla_x f(x)\| = 0.$$

Similarly, for $q = 2$, (2.7) is equivalent to

$$\|\nabla_x f(x)\| = 0 \quad \text{and} \quad \lambda_{\min}[\nabla_x^2 f(x)] \geq 0. \quad (2.9)$$

Its properties are further discussed in [9], but we emphasize that, for any $q \geq 1$ and in contrast with other known measures, it varies continuously when x varies continuously in \mathcal{F} . Solving the global optimization problem involved in its definition is easy for $q = 1$ as the global minimizer is analytically given by $d_* = -\delta \nabla_x^1 f(x) / \|\nabla_x^1 f(x)\|$, and also for $q = 2$ using a trust-region scheme (whose cost is essentially comparable to that of computing the leftmost eigenvalue in (2.9)). However this task may become NP-hard for larger q . This makes $\phi_{f,q}^\delta(x)$ an essentially theoretical tool for these cases. We note nevertheless that its calculation does not involve evaluating f or any of its derivatives, and its cost therefore does not affect the evaluation complexity of interest here.

If we now relax the notion of exact minimizers, we may define an (ϵ, δ) -approximate q th-order-necessary minimizer as a point $x \in \mathbb{R}^n$

$$\phi_{f,q}^\delta(x) \leq \epsilon \chi_q(\delta), \quad (2.10)$$

where

$$\chi_q(\delta) \stackrel{\text{def}}{=} \sum_{\ell=1}^q \frac{\delta^\ell}{\ell!} \quad (2.11)$$

provides a natural scaling. Again this notion reduces to familiar concepts in the low-order unconstrained cases. For instance, we verify that for unconstrained problems with $q = 2$, (2.10) requires that, if d is the global minimizer in (2.7) (the solution of a trust-region problem),

$$\max \left[0, -(\nabla_x^1 f(x)^T d + \tfrac{1}{2} d^T \nabla_x^2 f(x) d) \right] \leq \epsilon(\delta + \tfrac{1}{2} \delta^2),$$

which automatically holds for any $\delta \in (0, 1]$ if $\|\nabla_x^1 f(x)\| \leq \epsilon$ and $\lambda_{\min}[\nabla_x^2 f(x)] \geq -\epsilon$.

Having defined what we mean by high-order approximate minimizers, we now turn to describing what we mean by inaccurate objective function and derivatives values. It is important to observe at this point that, in an optimization problem, the role of the objective function is more central than that of any of its derivatives, since it is the quantity we ultimately wish to decrease. For this reason, we will handle the allowed inexactness in f differently from that in $\nabla_x^j f$: we will require an (adaptive) *absolute* accuracy for the first and a *relative* accuracy for the second. In fact, we can, in a first approach, abstract the relative accuracy requirements for the derivatives $\nabla_x^j f(x)$ into a requirement on the relative accuracy of $\Delta T_p^f(x, s)$. Let $\omega \in [0, 1]$ represent a relative accuracy level and denote inexact quantities with an overbar. For what follows, we will thus require that, if

$$\overline{\Delta T}_p^f(x, s, \omega) = \overline{T}_p^f(x_k, 0, \omega) - \overline{T}_p^f(x_k, s, \omega), \quad (2.12)$$

then

$$|\overline{\Delta T}_p^f(x, s, \omega) - \Delta T_p^f(x, s)| \leq \omega \overline{\Delta T}_p^f(x, s, \omega). \quad (2.13)$$

It may not be obvious at this point how to enforce this relative error bound: this is the object of Section 3 below. For now, we simply assume that it can be done in a finite number of evaluations of $\{\overline{\nabla}_x^j f(x)\}_{j=1}^p$ which are inexact approximations of $\{\nabla_x^j f(x)\}_{j=1}^p$.

Given an inexactly computed $\overline{\Delta T}_p^f(x, s, \omega)$ satisfying (2.13), we then have to consider to compute our optimality measure inexactly too. Observing that the definition (2.7) is independent of $f(x)$ because of cancellation, we see that

$$\overline{\phi}_{f,q}^\delta(x, \omega) = \max \left[0, \text{globmax}_{\substack{x+d \in \mathcal{F} \\ \|d\| \leq \delta}} \overline{\Delta T}_q^f(x, d, \omega) \right]. \quad (2.14)$$

Under the above assumptions, we now describe an algorithm allowing inexact computation of both the objective function and its derivatives whose purpose is to find (for given q and a suitable relative accuracy ω) a point x_k satisfying

$$\overline{\phi}_{f,q}^\delta(x, \omega) \leq \frac{\epsilon}{1 + \omega} \chi_q(\delta) \quad (2.15)$$

for some optimality radius $\delta \in (0, 1]$. This algorithm uses a regularized Taylor's model defined at iteration k by

$$m_k(s) \stackrel{\text{def}}{=} \overline{T}_p^f(x_k, s, \omega_k) + \frac{\sigma_k}{(p + \beta)!} \|s\|^{p+\beta}. \quad (2.16)$$

This model is then approximately minimized and the resulting trial point is then accepted or rejected depending on whether or not it produces a significant decrease. This is detailed in Algorithm 2.1 on the following page.

Algorithm 2.1: Adaptive Regularization of order p with Dynamic Accuracy (AR $_p$ DA)

Step 0: Initialization. An initial point $x_0 \in \mathcal{F}$ and an initial regularization parameter $\sigma_0 > 0$ are given, as well as an accuracy level $\epsilon \in (0, 1)$ and an initial relative accuracy $\omega_0 \geq 0$. The constants κ_ω , δ_{-1} , θ , μ , η_1 , η_2 , γ_1 , γ_2 , γ_3 and σ_{\min} are also given and satisfy $\theta > 0$, $\mu \in (0, 1]$, $\delta_{-1} \in (0, 1]$, $\sigma_{\min} \in (0, \sigma_0]$,

$$0 < \eta_1 \leq \eta_2 < 1, \quad 0 < \gamma_1 < 1 < \gamma_2 < \gamma_3, \quad (2.17)$$

$$\alpha \in (0, 1), \quad \kappa_\omega \in (0, \tfrac{1}{2}\alpha\eta_1] \quad \text{and} \quad \omega_0 = \min \left[\kappa_\omega, \frac{1}{\sigma_0} \right]. \quad (2.18)$$

Set $k = 0$.

Step 1: Compute the optimality measure and check for termination.

Compute $\bar{\phi}_{f,q}^{\delta_{k-1}}(x_k, \omega_k)$. If (2.15) holds with $\delta = \delta_{k-1}$, terminate with the approximate solution $x_\epsilon = x_k$.

Step 2: Step calculation. Attempt to compute a step $s_k \neq 0$ such that $x_k + s_k \in \mathcal{F}$ and an optimality radius $\delta_k \in (0, 1]$ by approximately minimizing the model $m_k(s)$ in the sense that

$$m_k(s_k) < m_k(0) \quad (2.19)$$

and

$$\|s_k\| \geq \mu \epsilon^{\frac{1}{p-q+\beta}} \quad \text{or} \quad \bar{\phi}_{m_k,q}^{\delta_k}(x_k + s_k, \omega_k) \leq \frac{\theta \|s_k\|^{p-q+\beta}}{(p-q+\beta)!} \chi_q(\delta_k). \quad (2.20)$$

If no such step exists, terminate with the approximate solution $x_\epsilon = x_k$.

Step 3: Acceptance of the trial point. Compute $\bar{f}_k(x_k + s_k, \omega_k)$ ensuring that

$$|\bar{f}_k(x_k + s_k, \omega_k) - f(x_k + s_k)| \leq \omega_k |\overline{\Delta T}_p^f(x_k, s_k, \omega_k)|. \quad (2.21)$$

Also ensure (by setting $\bar{f}_k(x_k, \omega_k) = \bar{f}_{k-1}(x_k, \omega_{k-1})$ or by (re)computing $\bar{f}_k(x_k, \omega_k)$) that

$$|\bar{f}_k(x_k, \omega_k) - f(x_k)| \leq \omega_k |\overline{\Delta T}_p^f(x_k, s_k, \omega_k)|. \quad (2.22)$$

Then define

$$\rho_k = \frac{\bar{f}_k(x_k, \omega_k) - \bar{f}_k(x_k + s_k, \omega_k)}{\overline{\Delta T}_p^f(x_k, s_k, \omega_k)}. \quad (2.23)$$

If $\rho_k \geq \eta_1$, then define $x_{k+1} = x_k + s_k$; otherwise define $x_{k+1} = x_k$.

Step 4: Regularization parameter update. Set

$$\sigma_{k+1} \in \begin{cases} [\max(\sigma_{\min}, \gamma_1 \sigma_k), \sigma_k] & \text{if } \rho_k \geq \eta_2, \\ [\sigma_k, \gamma_2 \sigma_k] & \text{if } \rho_k \in [\eta_1, \eta_2), \\ [\gamma_2 \sigma_k, \gamma_3 \sigma_k] & \text{if } \rho_k < \eta_1. \end{cases} \quad (2.24)$$

Step 5: Relative accuracy update. Set

$$\omega_{k+1} = \min \left[\kappa_\omega, \frac{1}{\sigma_{k+1}} \right]. \quad (2.25)$$

Increment k by one and go to Step 1.

Some comments on this algorithm are useful at this stage.

1. That Step 2 may not be able, for $q > 2$, to compute a nonzero step (and should then cause termination) can be seen by considering the following one-dimensional example. Let $p = q = 3$, $\mathcal{F} = \mathbb{R}$, $\omega_k = 0$ and $\delta_{k-1} = 1$ and suppose that $T_3(x_k, s) = s^2 - 2s^3$ and also that $\sigma_k = 24$. This implies that $m_k(s) = s^2 - 2s^3 + s^4 = s^2(1 - s)^2$ and we immediately see that the origin is a global minimizer of $m_k(s)$. But a simple calculation shows that $T_3(x_k, 1) = -1$ and hence termination will not occur in Step 1 if $\epsilon \leq 1/\chi_3(1) = \frac{4}{7}$. As a consequence, as was pointed out in [10], the possibility of a zero s_k cannot be ignored in Step 2, but will show below that this is acceptable.
2. Our assumption (2.13) is used three times in the algorithm: in Step 1 for computing $\overline{\phi}_{f,q}^{\delta_{k-1}}(x_k, \omega_k)$ and in Step 2 when computing s_k and $\overline{\phi}_{m_k,q}^{\delta_k}(x_k + s_k, \omega_k)$.
3. As indicated above, we require a bound on the absolute error in the objective function value: this is the object of (2.21) and (2.22), where we introduced the notation $\overline{f}_k(x_k, \omega_k)$ to denote an inexact approximation of $f(x_k)$. Note that a new value of $\overline{f}_k(x_k, \omega_k)$ should be computed to ensure (2.22) in Step 3 only if $k > 0$ and $\omega_{k-1} \overline{\Delta T}_p^f(x_{k-1}, s_{k-1}, \omega_{k-1}) > \omega_k \overline{\Delta T}_p^f(x_k, s_k, \omega_k)$. If this is the case the (inexact) function value is computed twice per iteration instead of just once.
4. At variance with the trust-region method with dynamic accuracy of [13, Section 10.6] and [1], we do not recompute approximate values of the objective function at x_k once the computation of s_k is complete (provided we can ensure (2.13), as discussed in Section 3).
5. If $\|s_k\| \geq \mu \epsilon^{\frac{1}{p-q+\beta}}$ in Step 2, then the (potentially costly) calculation of $\overline{\phi}_{m_k,q}^{\delta_k}(x_k + s_k, \omega_k)$ is unnecessary and δ_k may be chosen arbitrarily in $(0, 1]$.
6. We call iteration k *successful* when $\rho_k \geq \eta_1$ and $x_{k+1} = x_k + s_k$. The iteration is called *unsuccessful* otherwise, and $x_{k+1} = x_k$ in this case. We use the notation

$$\mathcal{S}_k = \{j \in \{0, \dots, k\} \mid \rho_j \geq \eta_1\} \quad (2.26)$$

to denote the set of successful iterations of index at most k .

7. As indicated above, ensuring (2.13) may require a certain number of (approximate) evaluations of the derivatives of f . For a single iteration of the algorithm, these evaluations are always at the current iterate x_k and we define τ_k to be a bound on the number of these evaluations at iterations k , that is

$$\begin{aligned} &\text{iteration } k \text{ of the ARpDA algorithm requires at most } \tau_k \\ &\text{approximate evaluations of } \overline{f} \text{ and } \{\nabla_x^j f(x_k)\}_{j=1}^p. \end{aligned} \quad (2.27)$$

We now state some properties of the algorithm that are derived without modification from the case where the computation of f and its derivatives are exact.

Lemma 2.2 [8, Theorem 2.1] The mechanism of the ARpDA algorithm ensures that, if

$$\sigma_k \leq \sigma_{\max}, \quad (2.28)$$

for some $\sigma_{\max} > 0$, then

$$k + 1 \leq |\mathcal{S}_k| \left(1 + \frac{|\log \gamma_1|}{\log \gamma_2} \right) + \frac{1}{\log \gamma_2} \log \left(\frac{\sigma_{\max}}{\sigma_0} \right). \quad (2.29)$$

This shows that the number of unsuccessful iterations must remain a fixed proportion of that of the successful ones.

Lemma 2.3 [9, Lemma 2.5] Suppose that $s_k^* \neq 0$ is a global minimizer of $m_k(s)$ under the constraint that $x_k + s \in \mathcal{F}$, such $m_k(s_k^*) < m_k(0)$. Then there exist a neighbourhood of s_k^* and a range of sufficiently small δ such that (2.19) and the second part of (2.20) hold for any s_k in the intersection of this neighbourhood with \mathcal{F} and any δ_k in this range.

This last lemma thus ensures that the algorithm is well-defined when $s_k \neq 0$. We discuss below what happens if a nonzero descent step cannot be computed.

3 Enforcing the relative error on Taylor increments

We now return to the question of enforcing (2.13) and deriving suitable bounds on τ_k . For improved readability, we temporarily ignore the iteration index k .

3.1 The accuracy checks

While there may be circumstances where (2.13) can be enforced directly, we consider here that the only control the user has on the accuracy of $\overline{\Delta T}_p^f(x, s, \omega)$ is by enforcing bounds $\{\varepsilon_j\}_{j=1}^p$ on the absolute errors on the derivative tensors $\{\nabla_x^j f(x)\}_{j=1}^p$. In other words, we seek to ensure (2.13) by selecting absolute accuracies $\{\varepsilon_j\}_{j=1}^p$ such that, when

$$\|\overline{\nabla_x^j f(x)} - \nabla_x^j f(x)\|_{[j]} \leq \varepsilon_j \quad \text{for } j \in \{1, \dots, p\}, \quad (3.1)$$

the desired accuracy requirement follows.

In all cases described below, the process can be viewed as an iteration with four main steps. The first is to compute the relevant approximate derivative satisfying (3.1) for given values of $\{\varepsilon_j\}_{j=1}^p$. The second is to use these approximate derivatives to compute the desired Taylor increment and associated quantities. Tests are then performed in the third step to verify the desired accuracy requirements and terminate if they are met. If not the case, the absolute accuracies $\{\varepsilon_j\}_{j=1}^p$ are then decreased before a new iteration is started.

As can be expected, a suitable relative accuracy requirement will be achievable when $\overline{\Delta T}_p^f(x, s, \omega)$ remains safely away from zero, but, if exact computations are to be avoided, we may have to accept a simpler absolute accuracy guarantee when $\overline{\Delta T}_p^f(x, s, \omega)$ vanishes.

Assume that a vector v_ω , a bound $\delta \geq \|v_\omega\|$, a degree r , the increment $\overline{\Delta T}_r(x, v_\omega, \omega)$, the current absolute accuracies $\{\varepsilon_j\}_{j=1}^p$, and the requested relative and absolute accuracies ω and $\xi > 0$ are given. (We use the notations $\overline{\Delta T}_r(x, v_\omega, \omega)$ and $\Delta T_r(x, v_\omega)$ without superscript because we intend to use the algorithm for $\overline{\Delta T}_q^f(x, v_\omega, \omega)$, $\overline{\Delta T}_p^f(x, v_\omega, \omega)$ and $\overline{\Delta T}_q^{m_k}(x, v_\omega, \omega)$.) We then formalize the resulting accuracy tests in the VERIFY algorithm, stated as Algorithm 3.1. Because it will be the case below, we assume for simplicity that $\overline{\Delta T}_r(x, v_\omega, \omega) \geq 0$.

Algorithm 3.1: Verify the accuracy of $\overline{\Delta T}_r(x, v_\omega, \omega)$

Set **flag** = 0.
 $\text{flag} = \text{VERIFY}\left(v_\omega, \delta, \overline{\Delta T}_r(x, v_\omega, \omega), \{\varepsilon_j\}_{j=1}^p, \omega, \xi\right)$

- If

$$\overline{\Delta T}_r(x, v_\omega, \omega) = 0 \quad \text{and} \quad \max_{j \in \{1, \dots, r\}} \varepsilon_j \leq \xi, \quad (3.2)$$

set **flag** = 1.

- Else, if

$$\overline{\Delta T}_r(x, v_\omega, \omega) > 0 \quad \text{and} \quad \sum_{j=1}^r \frac{\varepsilon_j}{j!} \delta^j \leq \omega \overline{\Delta T}_r(x, v_\omega, \omega), \quad (3.3)$$

set **flag** = 2.

- Else, if

$$\overline{\Delta T}_r(x, v_\omega, \omega) > 0 \quad \text{and} \quad \sum_{j=1}^r \frac{\varepsilon_j}{j!} \delta^j \leq \xi \chi_r(\delta), \quad (3.4)$$

set **flag** = 3.

Let us now consider what properties are ensured for the various possible values of **flag**.

Lemma 3.1 Assume (3.1) holds. Then we have that

- if
$$\max_{j \in \{1, \dots, p\}} \varepsilon_j \leq \xi, \quad (3.5)$$

then the VERIFY algorithm returns a nonzero **flag**,

- if the VERIFY algorithm terminates with **flag** = 1, then $\overline{\Delta T}_r(x, v_\omega, \omega) = 0$ and

$$|\overline{\Delta T}_r(x, v, \omega) - \Delta T_r(x, v)| \leq \xi \chi_r(\|v\|) \quad \text{for all } v, \quad (3.6)$$

- if the VERIFY algorithm terminates with **flag** = 2, then $\overline{\Delta T}_r(x, v_\omega, \omega) > 0$ and

$$|\overline{\Delta T}_r(x, v, \omega) - \Delta T_r(x, v)| \leq \omega \overline{\Delta T}_r(x, v_\omega, \omega), \quad \text{for all } v \text{ with } \|v\| \leq \delta, \quad (3.7)$$

- if the VERIFY algorithm terminates with **flag** = 3, then $\overline{\Delta T}_r(x, v_\omega, \omega) > 0$ and

$$\max [\overline{\Delta T}_r(x, v_\omega, \omega), |\overline{\Delta T}_r(x, v, \omega) - \Delta T_r(x, v)|] \leq \frac{\xi}{\omega} \chi_r(\delta) \quad \text{for all } v \text{ with } \|v\| \leq \delta. \quad (3.8)$$

Proof. We first prove the first proposition. If $\overline{\Delta T}_r(x, v_\omega, \omega) = 0$ and (3.5), then (3.2) ensures that **flag** = 1 is returned. If $\overline{\Delta T}_r(x, v_\omega, \omega) > 0$ we deduce that

$$\sum_{j=1}^r \frac{\varepsilon_j}{j!} \delta^j \leq \left[\max_{j \in \{1, \dots, r\}} \varepsilon_j \right] \chi_r(\delta) \leq \xi \chi_r(\delta)$$

also causing termination with **flag** = 3 because of (3.4) if it has not occurred with **flag** = 2 because of (3.3), hence proving the first proposition.

Consider now the three possible termination cases and suppose first that termination occurs with **flag** = 1. Then, using the triangle inequality, (3.1), (3.2) and (2.11), we have that, for any v ,

$$|\overline{\Delta T}_r(x, v, \omega) - \Delta T_r(x, v)| \leq \sum_{j=1}^r \frac{\varepsilon_j}{j!} \|v\|^j \leq \xi \chi_r(\|v\|)$$

yielding (3.6). Suppose now that **flag** = 2. Then $\overline{\Delta T}_r(x, v_\omega, \omega) > 0$, $v \neq 0$ and, for any v with $\|v\| \leq \delta$,

$$|\overline{\Delta T}_r(x, v, \omega) - \Delta T_r(x, v)| \leq \sum_{j=1}^r \frac{\varepsilon_j}{j!} \|v\|^j \leq \sum_{j=1}^r \frac{\varepsilon_j}{j!} \delta^j \leq \omega \overline{\Delta T}_r(x, v_\omega, \omega),$$

which is (3.7). Suppose finally that **flag** = 3. Since termination did not occur in (3.3), we have that

$$0 < \omega \overline{\Delta T}_r(x, v_\omega, \omega) \leq \xi \chi_r(\delta) \quad (3.9)$$

Hence, (3.4) implies that, for any v with $\|v\| \leq \delta$,

$$|\overline{\Delta T}_r(x, v, \omega) - \Delta T_r(x, v)| \leq \sum_{j=1}^r \frac{\varepsilon_j}{j!} \|v\|^j \leq \sum_{j=1}^r \frac{\varepsilon_j}{j!} \delta^j \leq \frac{\xi}{\omega} \chi_r(\delta).$$

This inequality and (3.9) together imply (3.8). \square

Clearly, the outcome corresponding to our initial aim to obtain a relative error at most ω corresponds to the case where **flag** = 2. As we will see below, the two other cases are also useful.

3.2 Computing $\overline{\phi}_{f,q}^{\delta_{k-1}}(x_k, \omega_k)$

We now consider, in Algorithm 3.2 on the next page, how to compute the optimality measure $\overline{\phi}_{f,q}^{\delta_{k-1}}(x_k, \omega_k)$ in Step 1 of the ARpDA algorithm.

We immediately observe that Algorithm 3.2 terminates in a finite number of iterations, since it does so as soon as **flag** > 0, which, because of the first proposition of Lemma 3.1, must happen after a finite number of passes in iterations using (3.10). We discuss in Section 3.4 exactly how many such decreases might be needed.

We now verify that terminating the ARpDA algorithm as indicated in this modified version of Step 1 provides the required result. We start noting that, if x_k is an isolated feasible point

Algorithm 3.2: Modified Step 1 of the AR_pDA algorithm**Step 1: Compute the optimality measure and check for termination.**

Step 1.0: The iterate x_k and the radius $\delta_{k-1} \in (0, 1]$ are given, as well as a constant $\gamma_\varepsilon \in (0, 1)$. Choose an initial set of derivative absolute accuracies $\{\varepsilon_{j,0}\}_{j=1}^p$ and set $i_\varepsilon = 0$.

Step 1.1: If unavailable, compute $\{\overline{\nabla}_{x^j}^j f(x_k)\}_{j=1}^q$ satisfying (3.1) with $\varepsilon_j = \varepsilon_{j,i_\varepsilon}$ for $j \in \{1, \dots, q\}$.

Step 1.2: Solve

$$\text{globmax}_{\substack{x_k + d \in \mathcal{F} \\ \|d\| \leq \delta_{k-1}}} \overline{\Delta T}_q^f(x_k, d, \omega_k),$$

to obtain the maximizer d_k and the corresponding Taylor increment $\overline{\Delta T}_q^f(x_k, d_k, \omega_k)$. Compute

$$\mathbf{flag} = \text{VERIFY}\left(d_k, \delta_{k-1}, \overline{\Delta T}_q^f(x_k, d_k, \omega_k), \{\varepsilon_j\}_{j=1}^p, \omega_k, \frac{1}{2}\omega_k \epsilon\right).$$

Step 1.3: Terminate the AR_pDA algorithm with the approximate solution $x_\epsilon = x_k$ if $\mathbf{flag} = 1$, or if $\mathbf{flag} = 3$, or if $\mathbf{flag} = 2$ and (2.15) holds with $\delta = \delta_{k-1}$. Also go to Step 2 of the AR_pDA algorithm if $\mathbf{flag} = 2$ but (2.15) fails.

Step 1.4: Otherwise (i.e. if $\mathbf{flag} = 0$), set

$$\varepsilon_{j,i_\varepsilon+1} = \gamma_\varepsilon \varepsilon_{j,i_\varepsilon} \quad \text{for } j \in \{1, \dots, p\}, \quad (3.10)$$

increment i_ε by one and return to Step 1.1.

(i.e. such that the intersection of any ball of radius $\delta_{k-1} > 0$ centered at x_k with \mathcal{F} is reduced to x_k), then clearly $d_k = 0$ and thus, irrespective of ω_k and $\delta_{k-1} > 0$,

$$\phi_{f,q}^{\delta_{k-1}}(x_k) = 0 = \overline{\Delta T}_q^f(x_k, d_k, \omega_k) = \overline{\phi}_{f,q}^{\delta_{k-1}}(x_k, \omega_k), \quad (3.11)$$

which means that $\overline{\phi}_{f,q}^{\delta_{k-1}}(x_k, \omega_k)$ is a faithful indicator of optimality at x_k .

Lemma 3.2 If the ARpDA algorithm terminates within Step 1.3, then

$$\phi_{f,q}^{\delta_{k-1}}(x_k) \leq \epsilon \chi_q(\delta_{k-1}) \quad (3.12)$$

and x_k is a (ϵ, δ_{k-1}) -approximate q th-order-necessary minimizer. Otherwise Algorithm 3.2 terminates with

$$(1 - \omega_k) \overline{\phi}_{f,q}^{\delta_{k-1}}(x_k, \omega_k) \leq \phi_{f,q}^{\delta_{k-1}}(x_k) \leq (1 + \omega_k) \overline{\phi}_{f,q}^{\delta_{k-1}}(x_k, \omega_k). \quad (3.13)$$

Proof. If x_k is an isolated feasible point, the lemma's conclusions directly follow from (3.11). Assume therefore that x_k is not an isolated feasible point and note first that, because Step 1.2 finds the global maximum of $\overline{\Delta T}_q^f(x_k, d, \omega_k)$, we have that $\overline{\Delta T}_q^f(x_k, d_k, \omega_k) \geq 0$. Suppose now that, in Step 1.2, the VERIFY algorithm returns **flag** = 1 and thus that $\overline{\Delta T}_q^f(x_k, d_k, \omega_k) = 0$. This means that x_k is a global minimizer of $\overline{T}_q^f(x_k, d, \omega_k)$ in the intersection of a ball of radius δ_{k-1} and \mathcal{F} and $\overline{\Delta T}_q^f(x_k, d, \omega_k) \leq 0$ for any d in this intersection. Thus, for any such d , we obtain from (3.6) with $\xi = \frac{1}{2}\omega_k\epsilon$ that

$$\Delta T_q^f(x_k, d) \leq \overline{\Delta T}_q^f(x_k, d, \omega_k) + \left| \overline{\Delta T}_q^f(x_k, d, \omega_k) - \Delta T_q^f(x_k, d) \right| \leq \frac{1}{2}\omega_k\epsilon\chi_q(\delta_{k-1}),$$

which, since $\omega_k \leq 1$, implies (3.12). Suppose next that the VERIFY algorithm returns **flag** = 3. Then $\overline{\Delta T}_q^f(x_k, d_k, \omega_k) > 0$ and thus $d_k \neq 0$. Using the fact that the nature of Step 1.2 ensures that $\overline{\Delta T}_q^f(x_k, d, \omega_k) \leq \overline{\Delta T}_q^f(x_k, d_k, \omega_k)$ for d with $\|d\| \leq \delta_{k-1}$ we have, using (3.8) with $\xi = \frac{1}{2}\omega_k\epsilon$, that, for all such d ,

$$\begin{aligned} \Delta T_q^f(x_k, d) &\leq \overline{\Delta T}_q^f(x_k, d, \omega_k) + \left| \overline{\Delta T}_q^f(x_k, d, \omega_k) - \Delta T_q^f(x_k, d) \right| \\ &\leq \overline{\Delta T}_q^f(x_k, d_k, \omega_k) + \left| \overline{\Delta T}_q^f(x_k, d, \omega_k) - \Delta T_q^f(x_k, d) \right| \\ &\leq \epsilon \chi_q(\delta_{k-1}) \end{aligned}$$

yielding (3.12). If the VERIFY algorithm returns **flag** = 2, then, for any d with $\|d\| \leq \delta_{k-1}$,

$$\Delta T_q^f(x_k, d) \leq \overline{\Delta T}_q^f(x_k, d, \omega_k) + \left| \overline{\Delta T}_q^f(x_k, d, \omega_k) - \Delta T_q^f(x_k, d) \right| \leq (1 + \omega_k) \overline{\Delta T}_q^f(x_k, d_k, \omega_k).$$

Thus, for all d with $\|d\| \leq \delta_{k-1}$,

$$\max \left[0, \Delta T_q^f(x_k, d) \right] \leq (1 + \omega_k) \max \left[0, \overline{\Delta T}_q^f(x_k, d_k, \omega_k) \right] = (1 + \omega_k) \overline{\phi}_{f,q}^{\delta_{k-1}}(x_k, \omega_k). \quad (3.14)$$

But termination implies that (2.15) holds for $\delta = \delta_{k-1}$, and (3.12) follows with this value of δ . Finally, if the ARpDA algorithm does not terminate within Step 1.3 but Algorithm 3.2 terminates, it must be because the VERIFY algorithm returns `flag` = 2. This implies, as above, that (3.14) holds, which is the rightmost part of (3.13). Similarly, for any d with $\|d\| \leq \delta_{k-1}$,

$$\begin{aligned} \Delta T_q^f(x_k, d) &\geq \overline{\Delta T}_q^f(x_k, d, \omega_k) - \left| \overline{\Delta T}_q^f(x_k, d, \omega_k) - \Delta T_q^f(x_k, d) \right| \\ &\geq \overline{\Delta T}_q^f(x_k, d, \omega_k) - \omega_k \overline{\Delta T}_q^f(x_k, d_k, \omega_k). \end{aligned}$$

Hence

$$\begin{aligned} \text{globmax}_{\substack{x_k+d \in \mathcal{F} \\ \|d\| \leq \delta_{k-1}}} \Delta T_q^f(x_k, d) &\geq \text{globmax}_{\substack{x_k+d \in \mathcal{F} \\ \|d\| \leq \delta_{k-1}}} \left[\overline{\Delta T}_q^f(x_k, d, \omega_k) - \omega_k \overline{\Delta T}_q^f(x_k, d_k, \omega_k) \right] \\ &= (1 - \omega_k) \overline{\Delta T}_q^f(x_k, d_k, \omega_k). \end{aligned}$$

Since $\overline{\Delta T}_q^f(x_k, d_k, \omega_k) > 0$ when the VERIFY algorithm returns `flag` = 2, we then obtain that, for all $\|d\| \leq \delta_{k-1}$,

$$\max_{\substack{x_k+d \in \mathcal{F} \\ \|d\| \leq \delta_{k-1}}} \left[0, \text{globmax} \Delta T_q^f(x_k, d) \right] \geq \max \left[0, (1 - \omega_k) \overline{\Delta T}_q^f(x_k, d_k, \omega_k) \right] = (1 - \omega_k) \overline{\phi}_{f,q}^{\delta_{k-1}}(x_k, \omega_k),$$

which is the leftmost part of (3.13). \square

3.3 Computing s_k

We now consider computing s_k at Step 2 of the ARpDA algorithm. The process is more complicated than for Step 1, as it potentially involves two situations in which one wishes to guarantee a suitable relative error. The first is when minimizing the model

$$m_k(s) = \overline{f}(x_k, \omega_k) - \overline{\Delta T}_p^f(x_k, s, \omega_k) + \frac{\sigma_k}{(p + \beta)!} \|s\|^{p+\beta}$$

or, equivalently, maximizing

$$-m_k(s) = -\overline{f}(x_k, \omega_k) + \overline{\Delta T}_p^f(x_k, s, \omega_k) - \frac{\sigma_k}{(p + \beta)!} \|s\|^{p+\beta}, \quad (3.15)$$

and the second is when globally minimizing the model's Taylor expansion taken at $x_k + s_k$ in a neighbourhood of diameter δ_k . The first of these situations can be handled in a way very similar to that used above for computing $\overline{\phi}_{f,q}^{\delta_{k-1}}(x_k)$ in Step 1: given a set of approximate derivatives, a step s_k is computed such that it satisfies (2.19) and (2.20), the relative error of the associated $\overline{\Delta T}_p^f(x_k, s_k, \omega_k)$ is then evaluated and, if it is insufficient, the accuracy on the derivative approximations improved and the process restarted. If the relative error on $\overline{\Delta T}_p^f(x_k, s_k, \omega_k)$ is satisfactory and the first test of (2.20) fails, it remains to check that the relative error on $\overline{\phi}_{m_k,q}^{\delta_k}(x_k + s_k, \omega_k)$ is also satisfactory. Moreover, as in the original ARpDA algorithm, we have to take into account the possibility that minimizing the model might result in a vanishing decrease. The resulting somewhat involved process is formalized in Algorithm 3.3 on the following page.

Algorithm 3.3: Modified Step 2 of the AR_pDA algorithm**Step 2: Step calculation.**

Step 2.0: The iterate x_k , the radius $\delta_{k-1} \in (0, 1]$, the constant $\gamma_\varepsilon \in (0, 1)$, the counter i_ε and the absolute accuracies $\{\varepsilon_{j,i_\varepsilon}\}_{j=1}^p$ are given.

Step 2.1: If unavailable, compute $\{\overline{\nabla_x^j f}(x_k)\}_{j=1}^p$ satisfying (3.1) with $\varepsilon_j = \varepsilon_{j,i_\varepsilon}$ for $j \in \{1, \dots, p\}$.

Step 2.2: Compute s_k such that $x_k + s_k \in \mathcal{F}$ by minimizing the model $m_k(s)$ to ensure that (2.19) and (2.20) hold for some $\delta_k \in (0, 1]$, in order to obtain s_k , $\overline{\Delta T}_p^f(x_k, s_k, \omega_k)$, $\delta_k \in (0, 1]$ and, if the first part of (2.20) fails, $d_k^{m_k}$ with $\|d_k^{m_k}\| \leq \delta_k$ and $\overline{\Delta T}_q^{m_k}(x_k + s_k, d_k^{m_k}, \omega_k)$. Compute

$$\text{flag}_s = \text{VERIFY}\left(s_k, \|s_k\|, \overline{\Delta T}_p^f(x_k, s_k, \omega_k), \{\epsilon_j\}_{j=1}^p, \omega_k, \tfrac{1}{2}\omega_k\epsilon\right).$$

If $\text{flag}_s = 0$ go to Step 2.5.

Step 2.3: If $\text{flag}_s = 1$ or $\text{flag}_s = 3$, compute

$$\underset{x_k + s \in \mathcal{F}}{\text{globmin}} m_k(s),$$

to obtain the minimizer s_k , $\overline{\Delta T}_p^f(x_k, s_k, \omega_k)$.

Set $d_k^{m_k} = 0 = \overline{\Delta T}_q^{m_k}(x_k + s_k, d_k^{m_k}, \omega_k)$ and compute

$$\text{flag}_s = \text{VERIFY}\left(s_k, \|s_k\|, \overline{\Delta T}_p^f(x_k, s_k, \omega_k), \{\epsilon_j\}_{j=1}^p, \omega_k, \tfrac{1}{2}\omega_k\epsilon\right).$$

If $\text{flag}_s = 0$ go to Step 2.5.

Step 2.4: If $\text{flag}_s = 1$ or $\text{flag}_s = 3$, terminate the AR_pDA algorithm with $x_\epsilon = x_k$. Otherwise, if $\|s_k\| \geq \mu\epsilon^{\frac{1}{p-q+\beta}}$ or if

$$\text{flag}_d = \text{VERIFY}\left(d_k^{m_k}, \delta_k, \overline{\Delta T}_q^{m_k}(x_k + s_k, d_k^{m_k}, \omega_k), \{\epsilon_j\}_{j=1}^p, \omega_k, \frac{\omega_k \theta \|s_k\|^{p-q+\beta}}{2(p-q+\beta)!}\right) > 0,$$

go to Step 3 of the AR_pDA algorithm with the step s_k , the associated $\overline{\Delta T}_p^f(x_k, s_k, \omega_k)$ and δ_k .

Step 2.5: Set (if $\text{flag}_s = 0$ or $\text{flag}_d = 0$),

$$\varepsilon_{j,i+1} = \gamma_\varepsilon \varepsilon_{j,i} \quad \text{for } j \in \{1, \dots, p\}, \quad (3.16)$$

increment i_ε by one and go to Step 2.1.

Again, Algorithm 3.3 must terminate in a finite number of iterations. Indeed, if $\mathbf{flag}_s = 1$ or $\mathbf{flag}_s = 3$ at the end of Step 2.2 or Step 2.3 after finitely many iterations, the conclusion is obvious. But, if $\mathbf{flag}_s = 2$ for all iterations, then the first proposition of Lemma 3.1 ensures that $\mathbf{flag}_d > 0$ after finitely many decreases in (3.16), also causing termination.

We note that, because of (3.15) and because Step 2.2 imposes (2.19), we have that $\overline{\Delta T}_p^f(x_k, s_k, \omega_k) \geq 0$ at the end of this step. Let us first consider the case where the calls to the VERIFY algorithm in Step 2.2 and in Step 2.3 both return $\mathbf{flag}_s = 1$ or $\mathbf{flag}_s = 3$. Since this implies that s_k is a global minimizer of m_k , we have that

$$\overline{\Delta T}_p^f(x_k, s, \omega_k) - \frac{\sigma_k}{(p+\beta)!} \|s\|^{p+\beta} \leq \overline{\Delta T}_p^f(x_k, s_k, \omega_k) - \frac{\sigma_k}{(p+\beta)!} \|s_k\|^{p+\beta}. \quad (3.17)$$

Thus, if $\|s\| \leq \|s_k\|$, then $\overline{\Delta T}_p^f(x_k, s, \omega_k) \leq \overline{\Delta T}_p^f(x_k, s_k, \omega_k)$. This implies that

$$\text{globmax}_{\substack{x_k+s \in \mathcal{F} \\ \|s\| \leq \|s_k\|}} \overline{\Delta T}_p^f(x_k, s, \omega_k) = \overline{\Delta T}_p^f(x_k, s_k, \omega_k).$$

We may now repeat the proof of Lemma 3.2 for the cases $\mathbf{flag} \in \{1, 3\}$, with q replaced by p and δ_{k-1} replaced by $\|s_k\|$, and deduce that

$$\phi_{f,p}^{\|s_k\|}(x_k) \leq \epsilon \chi_p(\|s_k\|). \quad (3.18)$$

As it turns out, this can be viewed as a stronger optimality condition than (2.10) since it implies that the p -th (rather than q -th with $q \leq p$) order Taylor expansion of f around x_k is bounded below by a correctly scaled multiple of ϵ , and in a possibly larger neighbourhood. It is thus more than reasonable to terminate the AR p DA algorithm with $x_\epsilon = x_k$, as stated at the beginning of Step 2.4 in the modified Step 2. Note that this replaces the possible termination outcome in Step 2 of the unmodified AR p DA algorithm.

Assume now that Algorithm 3.3 terminates in Step 2.4. This means that the VERIFY algorithm invoked in either Step 2.2 or Step 2.3 terminates with $\mathbf{flag}_s = 2$, and we deduce from (3.7) that

$$\left| \overline{\Delta T}_p^f(x_k, s_k, \omega_k) - \Delta T_p^f(x_k, s_k) \right| \leq \omega_k \overline{\Delta T}_p^f(x_k, s_k, \omega_k). \quad (3.19)$$

Moreover, the triangle inequality gives that, for all d with $\|d\| \leq \delta_k$,

$$\Delta T_q^{m_k}(x_k + s_k, d) \leq \overline{\Delta T}_q^{m_k}(x_k + s_k, d, \omega_k) + \left| \overline{\Delta T}_q^{m_k}(x_k + s_k, d, \omega_k) - \Delta T_q^{m_k}(x_k + s_k, d) \right|. \quad (3.20)$$

Assume furthermore that, in Step 2.4, Algorithm 3.3 terminates because $\mathbf{flag}_d = 1$ is returned by VERIFY. Then, using (3.20), the definition of $d_k^{m_k}$ and (3.6), we obtain that, for all d with $\|d\| \leq \delta_k$,

$$\begin{aligned} \Delta T_q^{m_k}(x_k + s_k, d) &\leq \overline{\Delta T}_q^{m_k}(x_k + s_k, d_k^{m_k}, \omega_k) + \left| \overline{\Delta T}_q^{m_k}(x_k + s_k, d, \omega_k) - \Delta T_q^{m_k}(x_k + s_k, d) \right| \\ &= \left| \overline{\Delta T}_q^{m_k}(x_k + s_k, d, \omega_k) - \Delta T_q^{m_k}(x_k + s_k, d) \right| \\ &\leq \frac{\theta \|s_k\|^{p-q+\beta}}{(p-q+\beta)!} \chi_q(\delta_k). \end{aligned} \quad (3.21)$$

If, instead, termination occurs with VERIFY returning $\mathbf{flag}_d = 2$, then (3.20), (3.7) and the bound $\omega_k \leq \kappa_\omega$ yield that, for all d with $\|d\| \leq \delta_k$,

$$\Delta T_q^{m_k}(x_k + s_k, d) \leq (1 + \omega_k) \overline{\Delta T}_q^{m_k}(x_k + s_k, d_k^{m_k}, \omega_k) \leq (1 + \kappa_\omega) \frac{\theta \|s_k\|^{p-q+\beta}}{(p-q+\beta)!} \chi_q(\delta_k). \quad (3.22)$$

Finally, if termination occurs instead because VERIFY returns $\mathbf{flag}_d = 3$, we deduce from the (3.20) and (3.8) that, for all d with $\|d\| \leq \delta_k$,

$$\Delta T_q^{m_k}(x_k + s_k, d) \leq \frac{\theta \|s_k\|^{p-q+\beta}}{(p-q+\beta)!} \chi_q(\delta_k) \quad (3.23)$$

Observe now that each of (3.21), (3.22) or (3.23) ensures that

$$\phi_{m_k, q}^{\delta_k}(x_k + s_k) \leq (1 + \omega_k) \frac{\theta \|s_k\|^{p-q+\beta}}{(p-q+\beta)!} \chi_q(\delta_k) \leq (1 + \kappa_\omega) \frac{\theta \|s_k\|^{p-q+\beta}}{(p-q+\beta)!} \chi_q(\delta_k). \quad (3.24)$$

We summarize this discussion in the following statement.

Lemma 3.3 Suppose that the modified Step 2 given on page 14 is used in the ARpDA algorithm. If this algorithm terminates within that step, then (3.18) holds. Otherwise we have that (2.19), (3.19) and either $\|s_k\| \geq \mu \epsilon^{\frac{1}{p-q+\beta}}$, or (3.24) hold.

3.4 The complexity of a single ARpDA iteration

The last part of this section is devoted to bounding the evaluation complexity of a single iteration of the ARpDA algorithm. The count in (approximate) objective function evaluations is the simplest: these only occur in Step 3 which requires at most two such evaluations.

Now observe that evaluations of $\{\nabla_x^j f\}_{j=1}^p$ possibly occur in Steps 1.2 and 2.2. However it is important to note that, within these steps, the *derivatives are evaluated only if the current values of the absolute errors are smaller than that used for the previous evaluations* of the same derivative at the same point (x_k) . Moreover, these absolute errors are, by construction, linearly decreasing with rate γ_ϵ within the same iteration of the ARpDA algorithm (they are initialized in Step 1.0, decreased each time by a factor γ_ϵ in (3.10) invoked in Step 1.4, down to values $\{\epsilon_{j, i_\epsilon}\}_{j=1}^p$ which are then passed to the modified Step 2, and decreased there further in (3.16) in Step 2.5, again by successive multiplication with γ_ϵ). Furthermore, we have argued already, both for the modified Step 1 and the modified Step 2, that any of these algorithms terminates as soon as (3.5) holds for the relevant value of ξ , which we therefore need to determine. For Step 1, this value is $\frac{1}{2}\omega_k\epsilon$, while, for Step 2, one might think, at first sight, that it is given by

$$\min \left[\frac{1}{2}\omega_k\epsilon, \frac{\omega_k\theta\|s_k\|^{p-q+\beta}}{2(p-q+\beta)!} \right]. \quad (3.25)$$

There is however a slight complication, because the derivatives which are approximated in the VERIFY call of Step 2.4 are not $\{\nabla_x^j f(x_k)\}_{j=1}^q$, but $\{\nabla_d^j T_q^{m_k}(x_k + s_k, 0)\}_{j=1}^q$. It is easy

to verify that these (approximate) derivatives are given by

$$\overline{\nabla_d^j T_q^{m_k}}(x_k + s_k, 0) = \sum_{\ell=j}^p \frac{\overline{\nabla_x^\ell f(x_k)} s_k^{\ell-j}}{(\ell-j)!} + \left[\nabla_s^j \|s\|^{p+\beta} \right]_{s=s_k}, \quad (3.26)$$

where the last term of the right-hand side is exact. This yields the following error bound.

Lemma 3.4 Assume $\|s_k\| \leq \mu \epsilon^{\frac{1}{p-q+\beta}}$. Then, for all $j \in \{1, \dots, p\}$,

$$\left| \overline{\nabla_d^j T_q^{m_k}}(x_k + s_k, 0) - \nabla_d^j T_q^{m_k}(x_k + s_k, 0) \right| \leq 3 \varepsilon_j. \quad (3.27)$$

Proof. Using the triangle inequality, (3.26), the inequality $\|s_k\| \leq \mu \epsilon^{\frac{1}{p-q+\beta}} \leq \mu$ and (2.11), we have that, for all $j \in \{1, \dots, p\}$,

$$\left| \overline{\nabla_d^j T_q^{m_k}}(x_k + s_k, 0) - \nabla_d^j T_q^{m_k}(x_k + s_k, 0) \right| \leq \sum_{\ell=j}^p \frac{\varepsilon_j s_k^{\ell-j}}{(\ell-j)!} \leq \varepsilon_j \sum_{\ell=j}^p \frac{\mu^{\ell-j}}{(\ell-j)!} \leq \varepsilon_j (1 + \chi_p(\mu))$$

and (3.27) follows since $\chi_p(\mu) \leq 2\mu$. \square

This means that, in Step 2.4 when $\|s_k\| \leq \mu \epsilon^{\frac{1}{p-q+\beta}}$, we have to make $\{\overline{\nabla_x^j f(x_k)}\}_{j=1}^p$ three times more accurate than the desired accuracy in $\{\overline{\nabla_d^j T_q^{m_k}}(x_k + s_k, 0)\}_{j=1}^q$. In other words, if we focus on evaluations of $\{\overline{\nabla_x^j f(x_k)}\}_{j=1}^p$ rather than of $\{\overline{\nabla_d^j T_q^{m_k}}(x_k + s_k, 0)\}_{j=1}^q$, the relevant value for ξ in Step 2.4 should be divided by three, and become

$$\frac{\omega_k \theta \|s_k\|^{p-q+\beta}}{6(p-q+\beta)!} \quad \text{instead of} \quad \frac{\omega_k \theta \|s_k\|^{p-q+\beta}}{2(p-q+\beta)!}.$$

We may thus conclude from Lemma 3.1 that no further reduction in $\{\varepsilon_j\}_{j=1}^p$ (and hence no further approximation of $\{\overline{\nabla_x^j f(x_k)}\}_{j=1}^p$) will occur once i_ε , the number of decreases in $\{\varepsilon_j\}_{j=1}^p$, is large enough to ensure that

$$\gamma_\varepsilon^i \left[\max_{j \in \{1, \dots, p\}} \varepsilon_{j,0} \right] \leq \min \left[\frac{\omega_k \epsilon}{2}, \frac{\omega_k \theta \mu^{p-q+\beta}}{6(p-q+\beta)!}, \max_{j \in \{1, \dots, p\}} \varepsilon_{j,0} \right],$$

where the last term reflects the possibility that the initial $\{\varepsilon_{j,0}\}_{j=1}^p$ may already be small enough. Denoting $\kappa_\varepsilon = \max_{j \in \{1, \dots, p\}} \varepsilon_{j,0}$, we therefore deduce the following bound.

Lemma 3.5 Suppose that $\omega_k \geq \omega_{\min} > 0$ for all k . Then each iteration of the ARpDA algorithm involves at most 2 (approximate) evaluations of the objective function and at most

$$\left\lceil \frac{1}{\log(\gamma_\varepsilon)} \left\{ \log \left(\min \left[\frac{\omega_{\min} \epsilon}{2}, \frac{\omega_{\min} \theta \mu^{p-q+\beta}}{6(p-q+\beta)!}, \kappa_\varepsilon \right] \right) - \log(\kappa_\varepsilon) \right\} \right\rceil$$

(approximate) evaluations of its p first derivatives.

Note that, for simplicity, we have ignored the fact that only $q \leq p$ derivatives need to be evaluated in Steps 1.1. Lemma 3.5 can obviously be refined to reflect this observation. Globalizing instead the evaluations of \bar{f} and those of $\{\nabla_x^j f\}_{j=1}^p$, this lemma also implies that τ_k (as defined in (2.27)) satisfies the bound

$$\tau_k \leq \tau_{\max}(\epsilon) \stackrel{\text{def}}{=} \max \left\{ 2, \left\lfloor \frac{1}{\log(\gamma_\epsilon)} \left\{ \log \left(\min \left[\frac{\omega_{\min} \epsilon}{2}, \frac{\omega_{\min} \theta \mu^{p-q+\beta}}{6(p-q+\beta)!}, \kappa_\epsilon \right] \right) - \log(\kappa_\epsilon) \right\} \right\rfloor \right\}. \quad (3.28)$$

We conclude this section by a comment on what happens whenever exact derivatives are used, that is when $\kappa_\epsilon = 0$. In that case the (exact) derivatives are computed only once per iteration of the ARpDA algorithm (in Step 1.2 for the first q and in Step 2.2 for the remaining $p - q$) and every other call to VERIFY returns **flag** = 1 or **flag** = 2. Moreover, there is no need to recompute \bar{f} to obtain (2.22) in Step 3. The evaluation complexity of a single iteration of the ARpDA algorithm then reduces to a single evaluation of f and its first p derivatives (and $\tau_k = 1$ for all k), as expected.

4 Evaluation complexity of the deterministic ARpDA

This section is devoted to the evaluation complexity analysis of the ARpDA algorithm in the deterministic context. We start by providing a simple lower bound on the model decrease.

Lemma 4.1 [9, Lemma 3.1] The mechanism of the ARpDA algorithm guarantees that, for all $k \geq 0$,

$$\overline{\Delta T}_p^f(x_k, s_k, \omega_k) > \frac{\sigma_k}{(p + \beta)!} \|s_k\|^{p+\beta}, \quad (4.1)$$

and so (2.23) is well-defined.

Proof. We have that

$$0 < m_k(0) - m_k(s_k) = \overline{T}_p(x_k, 0, \omega_k) - \overline{T}_p(x_k, s_k, \omega_k) - \frac{\sigma_k}{(p + \beta)!} \|s_k\|^{p+\beta}.$$

□

We next show that the regularization parameter σ_k has to remain bounded, even in the presence of inexact computation of f and its derivatives. This lemma hinges heavily on (2.13), (2.21) and (2.22).

Lemma 4.2 Let $f \in \mathcal{C}^{p,\beta}(\mathbb{R}^n)$. Then, for all $k \geq 0$,

$$\sigma_k \leq \sigma_{\max} \stackrel{\text{def}}{=} \max \left[\sigma_0, \frac{\gamma_3(L+3)}{1 - \eta_2} \right] \quad (4.2)$$

and

$$\omega_k \geq \omega_{\min} \stackrel{\text{def}}{=} \min \left[\kappa_\omega, \frac{1}{\sigma_{\max}} \right]. \quad (4.3)$$

Proof. Assume that

$$\sigma_k \geq \frac{L+3}{1-\eta_2}. \quad (4.4)$$

Also observe that, because of the triangle inequality, (3.19) (as ensured by Lemma 3.3) and (2.22),

$$\begin{aligned} |\overline{T}_p^f(x_k, s_k, \omega_k) - T_p^f(x_k, s_k)| &\leq |\overline{f}_k(x_k, \omega_k) - f(x_k)| \\ &\quad + |\overline{\Delta T}_p^f(x_k, s_k, \omega_k) - \Delta T_p^f(x_k, s_k)| \\ &\leq 2\omega_k |\overline{\Delta T}_p^f(x_k, s_k, \omega_k)| \end{aligned}$$

and hence, again using the triangle inequality, (2.21), (2.5), (2.25), (4.1) and (4.4),

$$\begin{aligned} |\rho_k - 1| &\leq \frac{|\overline{f}_k(x_k + s_k, \omega_k) - \overline{T}_p^f(x_k, s_k, \omega_k)|}{\overline{\Delta T}_p^f(x_k, s_k, \omega_k)} \\ &\leq \frac{1}{\overline{\Delta T}_p^f(x_k, s_k, \omega_k)} \left[|\overline{f}_k(x_k + s_k, \omega_k) - f(x_k + s_k)| + |f(x_k + s_k) - T_p^f(x_k, s_k)| \right. \\ &\quad \left. + |\overline{T}_p^f(x_k, s_k, \omega_k) - T_p^f(x_k, s_k)| \right] \\ &\leq \frac{1}{\overline{\Delta T}_p^f(x_k, s_k, \omega_k)} \left[|f(x_k + s_k) - T_p^f(x_k, s_k)| + 3\omega_k \overline{\Delta T}_p^f(x_k, s_k, \omega_k) \right] \\ &\leq \frac{1}{\overline{\Delta T}_p^f(x_k, s_k, \omega_k)} \left[\frac{L}{(p+\beta)!} \|s_k\|^{p+\beta} + \frac{3\overline{\Delta T}_p^f(x_k, s_k, \omega_k)}{\sigma_k} \right] \\ &< \frac{L}{\sigma_k} + \frac{3}{\sigma_k} \\ &< 1 - \eta_2 \end{aligned}$$

and thus that $\rho_k > \eta_2$. Then iteration k is very successful in that $\rho_k \geq \eta_2$ and, because of (2.24), $\sigma_{k+1} \leq \sigma_k$. As a consequence, the mechanism of the algorithm ensures that (4.2) holds. Observe now this result and (2.25) imply that, for all k , ω_k may be chosen such that $\min[\kappa_\omega, \sigma_{\max}^{-1}] \leq \omega_k \leq \kappa_\omega$, yielding (4.3). \square

It is important to note that (4.3) in this lemma provides the lower bound on ω_k required in Lemma 3.5. We now borrow a technical result from [9].

Lemma 4.3 [9, Lemma 2.4] Let s be a vector of \mathbb{R}^n and $p \in \mathbb{N}_0$ and $\beta \in (0, 1]$ such that $j \in \{0, \dots, p\}$. Then

$$\|\nabla_s^j(\|s\|^{p+\beta})\|_{[j]} \leq \frac{(p+\beta)!}{(p-j+\beta)!} \|s\|^{p-j+\beta}. \quad (4.5)$$

Our next move is to prove a lower bound on the step norm. While the proof of this result is clearly inspired from that of [9, Lemma 3.3], it nevertheless crucially differs when approximate values are considered instead of exact ones.

Lemma 4.4 Let $f \in \mathcal{C}^{p,\beta}(\mathbb{R}^n)$. Then, for all $k \geq 0$ such that the ARpDA algorithm does not terminate at iteration $k + 1$,

$$\|s_k\| \geq \kappa_s \epsilon^{\frac{1}{p-q+\beta}}, \quad (4.6)$$

where

$$\kappa_s \stackrel{\text{def}}{=} \min \left\{ \mu, \left[\frac{(1 - \kappa_\omega)(p - q + \beta)!}{(1 + \kappa_\omega)(L + \sigma_{\max} + \theta(1 + \kappa_\omega))} \right]^{\frac{1}{p-q+\beta}} \right\}. \quad (4.7)$$

Proof. If $\|s_k\| > \mu \epsilon^{\frac{1}{p-q+\beta}}$, the result is obvious. Suppose now that $\|s_k\| \leq \mu \epsilon^{\frac{1}{p-q+\beta}}$. Since the algorithm does not terminate at iteration $k + 1$, we have that

$$\overline{\phi}_{f,q}^{\delta_k}(x_{k+1}) > \frac{\epsilon}{1 + \omega_k} \chi_q(\delta_k)$$

and therefore, using (3.13), that

$$\phi_{f,q}^{\delta_k}(x_{k+1}) > \frac{1 - \omega_k}{1 + \omega_k} \epsilon \chi_q(\delta_k). \quad (4.8)$$

Let the global minimum in the definition of $\phi_{f,q}^{\delta_k}(x_{k+1})$ be achieved at d with $\|d\| \leq \delta_k$. Then, using that $\|s_k\| \leq \mu \epsilon^{\frac{1}{p-q+\beta}}$ and successively using (2.7) (both for f at x_{k+1} and for m_k at s_k), the triangle inequality and (4.5), we deduce that

$$\begin{aligned} \phi_{f,q}^{\delta_k}(x_{k+1}) &= - \sum_{\ell=1}^q \frac{1}{\ell!} \nabla_x^\ell f(x_{k+1}) [d]^\ell \\ &\leq \left| \sum_{\ell=1}^q \frac{1}{\ell!} \nabla_x^\ell f(x_{k+1}) [d]^\ell - \sum_{\ell=1}^q \frac{1}{\ell!} \nabla_s^\ell T_p(x_k, s_k) [d]^\ell \right| - \sum_{\ell=1}^q \frac{1}{\ell!} \nabla_s^\ell T_p(x_k, s_k) [d]^\ell \\ &\leq \sum_{\ell=1}^q \frac{1}{\ell!} \left[\|\nabla_x^\ell f(x_{k+1}) - \nabla_s^\ell T_p(x_k, s_k)\|_{[\ell]} \right] \delta_k^\ell \\ &\quad - \sum_{\ell=1}^q \frac{1}{\ell!} \left(\nabla_s^\ell \left[T_p(x_k, s) + \frac{\sigma_k}{(p + \beta)!} \|s\|^{p+\beta} \right]_{s=s_k} \right) [d]^\ell \\ &\quad + \sum_{\ell=1}^q \frac{\sigma_k}{\ell! (p - \ell + \beta)!} \|s_k\|^{p-\ell+\beta} \delta_k^\ell. \end{aligned} \quad (4.9)$$

Now, because of (2.16), (2.14) and the fact that $\|d\| \leq \delta_k$, we have that

$$- \sum_{\ell=1}^q \frac{1}{\ell!} \left(\nabla_s^\ell \left[T_p(x_k, s) + \frac{\sigma_k}{(p + \beta)!} \|s\|^{p+\beta} \right]_{s=s_k} \right) [d]^\ell = \Delta T_q^{m_k}(x_k + s_k, d) \leq \phi_{m_k,q}^{\delta_k}(x_k + s_k)$$

and we may then use (3.24) (ensured by Lemma 3.3) to deduce from (4.9) that

$$\begin{aligned} \phi_{f,q}^{\delta_k}(x_{k+1}) &\leq \sum_{\ell=1}^q \frac{L}{\ell!(p-\ell+\beta)!} \|s_k\|^{p-\ell+\beta} \delta_k^\ell + (1+\kappa_\omega) \frac{\theta \chi_q(\delta_k)}{(p-q+\beta)!} \|s_k\|^{p-q+\beta} \\ &\quad + \sum_{\ell=1}^q \frac{\sigma_k}{\ell!(p-\ell+\beta)!} \|s_k\|^{p-\ell+\beta} \delta_k^\ell \\ &\leq \frac{[L + \sigma_k + \theta(1+\kappa_\omega)] \chi_q(\delta_k)}{(p-q+\beta)!} \|s_k\|^{p-q+\beta}, \end{aligned} \tag{4.10}$$

where our assumption that $\|s_k\| \leq \mu \epsilon^{\frac{1}{p-q+\beta}} \leq 1$ is used to obtain the last inequality. Using (4.8) and the bound $\omega_k \leq \kappa_\omega$, we thus have that

$$\|s_k\| \geq \min \left\{ \mu \epsilon^{\frac{1}{p-q+\beta}}, \left[\frac{\epsilon(1-\kappa_\omega)(p-q+\beta)!}{(1+\kappa_\omega)(L + \sigma_k + \theta(1+\kappa_\omega))} \right]^{\frac{1}{p-q+\beta}} \right\},$$

and (4.6) immediately follows from (4.2). \square

We now combine all the above results to deduce an upper bound on the maximum number of successful iterations, from which a final complexity bound immediately follows.

Theorem 4.5 Let $f \in \mathcal{C}^{p,\beta}(\mathbb{R}^n)$. Then, given $\epsilon \in (0,1)$, the ARpDA algorithm using the modified Steps 1 (on page 11) and 2 (on page 14) needs at most

$$\left\lceil \kappa_p(f(x_0) - f_{\text{low}}) \left(\epsilon^{-\frac{p+\beta}{p-q+\beta}} \right) \right\rceil + 1 \tag{4.11}$$

successful iterations (each involving at most $\tau_{\max}(\epsilon)$ evaluations of f and its p first derivatives) and at most

$$\left\lceil \tau_{\max}(\epsilon) \left\{ \left\lceil \kappa_p(f(x_0) - f_{\text{low}}) \left(\epsilon^{-\frac{p+\beta}{p-q+\beta}} \right) \right\rceil + 1 \right\} \left(1 + \frac{|\log \gamma_1|}{\log \gamma_2} \right) + \frac{1}{\log \gamma_2} \log \left(\frac{\sigma_{\max}}{\sigma_0} \right) \right\rceil \tag{4.12}$$

approximate evaluations of f and its derivatives in total to produce an iterate x_ϵ such that (2.10) or (3.18) holds, where σ_{\max} is given by (4.2), ω_{\min} by (4.3), $\tau_{\max}(\epsilon)$ by (3.28), and where

$$\kappa_p \stackrel{\text{def}}{=} \frac{(p+\beta)!}{\eta_1(1-\alpha)\sigma_{\min}} \max \left\{ \frac{1}{\mu^{p+\beta}}, \left[\frac{(1+\kappa_\omega)(L + \sigma_{\max} + \theta(1+\kappa_\omega))}{(1-\kappa_\omega)(p-q+\beta)!} \right]^{\frac{p+\beta}{p-q+\beta}} \right\}.$$

Proof. At each successful iteration k before termination the algorithm guarantees the

decrease

$$\begin{aligned}
f(x_k) - f(x_{k+1}) &\geq [\bar{f}_k(x_k, \omega_k) - \bar{f}_k(x_{k+1}, \omega_k)] - 2\omega_k \overline{\Delta T}_p^f(x_k, s_k, \omega_k) \\
&\geq \eta_1 \overline{\Delta T}_p^f(x_k, s_k, \omega_k) - \alpha \eta_1 \overline{\Delta T}_p^f(x_k, s_k, \omega_k) \\
&\geq \frac{\eta_1(1-\alpha)\sigma_{\min}}{(p+\beta)!} \|s_k\|^{p+\beta},
\end{aligned} \tag{4.13}$$

where we used (2.18), (2.23), (4.1) and (2.24). Moreover we deduce from (4.13), (4.6) and (4.2) that

$$f(x_k) - f(x_{k+1}) \geq \kappa_p^{-1} \epsilon^{\frac{p+\beta}{p-q+\beta}} \quad \text{where} \quad \kappa_p^{-1} \stackrel{\text{def}}{=} \frac{\eta_1(1-\alpha)\sigma_{\min}\kappa_s^{p+\beta}}{(p+\beta)!}. \tag{4.14}$$

Thus, since $\{f(x_k)\}$ decreases monotonically,

$$f(x_0) - f(x_{k+1}) \geq \kappa_p^{-1} \epsilon^{\frac{p+\beta}{p-q+\beta}} |\mathcal{S}_k|.$$

Using that f is bounded below by f_{low} , we conclude that

$$|\mathcal{S}_k| \leq \kappa_p(f(x_0) - f_{\text{low}}) \epsilon^{-\frac{p+\beta}{p-q+\beta}} \tag{4.15}$$

until termination, and the desired bound on the number of successful iterations follows. Lemma 2.2 is then invoked to compute the upper bound on the total number of iterations. \square

We emphasize that (4.11) was shown in [9] to be optimal for a quite wide class of minimization algorithms. The slightly weaker (4.12) may be seen as the (very modest) price to pay for allowing inexact evaluations.

Focusing on the order in ϵ and using (3.28), we therefore obtain what is, to the author's knowledge, the most complete result so far on evaluation complexity for nonconvex optimization.

Theorem 4.6 Let $f \in \mathcal{C}^{p,\beta}(\mathbb{R}^n)$. Then, given $\epsilon \in (0, 1)$, the ARpDA algorithm using the modified Steps 1 (on page 11) and 2 (on page 14) needs at most

$$O\left(\epsilon^{-\frac{p+\beta}{p-q+\beta}}\right) \text{ iterations}$$

and at most

$$O\left(|\log(\epsilon)| \epsilon^{-\frac{p+\beta}{p-q+\beta}}\right) \text{ (approximate) evaluations of } f \text{ and its } p \text{ first derivatives}$$

to compute an (ϵ, δ) -approximate q -th-order-necessary minimizer for the set-constrained problem (2.1).

In particular, if the p -th derivative of f is assumed to be globally Lipschitz rather than merely Hölder continuous (i.e. if $\beta = 1$), these order reduce to

$$O\left(\epsilon^{-\frac{p+1}{p-q+1}}\right) \text{ iterations}$$

and at most

$$O\left(|\log(\epsilon)|\epsilon^{-\frac{p+1}{p-q+1}}\right) \quad (\text{approximate}) \text{ evaluations of } f \text{ and its } p \text{ first derivatives.}$$

As indicated in the comment at the end of Section 3, all $O(|\log(\epsilon)|)$ terms reduce to a constant independent of ϵ if exact evaluations of f and its derivatives are used, and the above results then recover the optimal complexity results of [9].

5 Application to unconstrained and bound-constrained first- and second-order nonconvex inexact optimization

Because of its wide-ranging applicability, the framework discussed above may appear somewhat daunting in its generality. Moreover, the fact that it involves (possibly constrained) global optimization subproblems in several of its steps may suggest that it has to remain conceptual. We show in this section that *this is not the case*, and stress that it is *much simpler when specialized to small values of p and q* (which are, for now, the most practical ones) and that our approach leads to elegant and implementable numerical algorithms. To illustrate this point, we now review what happens for $p \leq 2$.

We first discuss the case where one seeks to compute a first-order critical point for an unconstrained optimization problems using approximate function values as well as approximate first derivatives. For simplicity of exposition, we will also assume that the gradient of f is Lipschitz (rather than Hölder) continuous. In our general context, this means that we consider the case where $q = 1$, $p = 1$, $\beta = 1$ and $\mathcal{F} = \mathbb{R}^n$. We first note that, as pointed out in (2.8),

$$\phi_{f,1}^\delta(x) = \|\nabla_x^1 f(x)\|\delta \quad \text{and} \quad \overline{\phi}_{f,1}^\delta = \|\overline{\nabla_x^1 f(x)}\|\delta \quad \text{irrespective of } \delta \in (0, 1], \quad (5.1)$$

which means that, since we can choose $\delta = 1$, Step 1 of the ARpdA algorithm reduces to the computation of an approximate gradient $\overline{\nabla_x^1 f}(x_k)$ with relative error ω_k and verification that ϵ -approximate optimality is not yet achieved. If that is the case, computing s_k at Step 2 is also extremely simple since it easy to verify that

$$s_k = s_k^* = -\frac{1}{\sigma_k} \overline{\nabla_x^1 f}(x_k).$$

Lemma 2.3 then ensures that this step is acceptable for some $\delta_k \in (0, 1]$, the value of which being irrelevant since it is not used in Step 1 of the next iteration. Moreover, if the relative error on $\overline{\nabla_x^1 f}(x_k)$ is bounded by ω_k , then

$$\begin{aligned} |\overline{\Delta T}_1^f(x_k, s_k) - \Delta T_1^f(x_k, s_k)| &\leq \|\overline{\nabla_x^1 f}(x_k) - \nabla_x^1 f(x_k)\| \frac{\|\overline{\nabla_x^1 f}(x_k)\|}{\sigma_k} \\ &\leq \omega_k \frac{\|\overline{\nabla_x^1 f}(x_k)\|^2}{\sigma_k} \\ &= \omega_k \overline{\Delta T}_1^f(x_k, s_k) \end{aligned}$$

and (2.13) automatically holds, so that no iteration is needed in Algorithm 3.3. The resulting algorithm, where we have made the modified Step 1 explicit, is given as Algorithm 5.1 (AR1DA) on the following page.

Algorithm 5.1: The AR1DA Algorithm

Step 0: Initialization. An initial point $x_0 \in \mathcal{F}$ and an initial regularization parameter $\sigma_0 > 0$ are given, as well as an accuracy level $\epsilon \in (0, 1)$ and an initial relative accuracy $\omega_0 \geq 0$. The constants $\alpha, \kappa_\omega, \kappa_\epsilon, \eta_1, \eta_2, \gamma_1, \gamma_2, \gamma_3$ and σ_{\min} are also given and satisfy $\sigma_{\min} \in (0, \sigma_0]$,

$$0 < \eta_1 \leq \eta_2 < 1, \quad 0 < \gamma_1 < 1 < \gamma_2 < \gamma_3, \\ \kappa_\epsilon \in (0, 1] \quad \alpha \in (0, 1), \quad \kappa_\omega \in (0, \tfrac{1}{2}\alpha\eta_1] \quad \text{and} \quad \omega_0 = \min \left[\kappa_\omega, \frac{1}{\sigma_0} \right].$$

Set $k = 0$.

Step 1: Compute the optimality measure and check for termination.

Initialize $\varepsilon_{1,0} = \kappa_\epsilon$ and set $i = 0$. Do

1. compute $\overline{\nabla_x^1 f}(x_k)$ with $\|\overline{\nabla_x^1 f}(x_k) - \nabla_x^1 f(x_k)\| \leq \varepsilon_{1,i}$ and increment i by one.
2. if $\|\overline{\nabla_x^1 f}(x_k)\| \leq \epsilon/(2(1 + \omega_k))$, terminate with $x_\epsilon = x_k$;
3. if $\varepsilon_{1,i} \leq \omega_k \|\overline{\nabla_x^1 f}(x_k)\|$, go to Step 2;
4. set $\varepsilon_{1,i+1} = \gamma_\epsilon \varepsilon_{1,i}$ and return to item 1 in this enumeration.

Step 2: Step calculation. Set

$$s_k = -\overline{\nabla_x^1 f}(x_k)/\sigma_k \quad \text{and} \quad \overline{\Delta T_p^f}(x_k, s_k, \omega_k) = \|\overline{\nabla_x^1 f}(x_k)\|^2/\sigma_k.$$

Step 3: Acceptance of the trial point.

Compute $\overline{f}_k(x_k + s_k, \omega_k)$ ensuring that

$$|\overline{f}_k(x_k + s_k, \omega_k) - f(x_k + s_k)| \leq \omega_k |\overline{\Delta T_p^f}(x_k, s_k, \omega_k)|. \quad (5.2)$$

Also ensure (by setting $\overline{f}_k(x_k, \omega_k) = \overline{f}_{k-1}(x_k, \omega_{k-1})$ or by (re)computing $\overline{f}_k(x_k, \omega_k)$) that

$$|\overline{f}_k(x_k, \omega_k) - f(x_k)| \leq \omega_k |\overline{\Delta T_p^f}(x_k, s_k, \omega_k)| \quad (5.3)$$

Then define

$$\rho_k = \frac{\overline{f}_k(x_k, \omega_k) - \overline{f}_k(x_k + s_k, \omega_k)}{\overline{\Delta T_p^f}(x_k, s_k, \omega_k)}. \quad (5.4)$$

If $\rho_k \geq \eta_1$, then define $x_{k+1} = x_k + s_k$; otherwise define $x_{k+1} = x_k$.

Step 4: Regularization parameter update. Set

$$\sigma_{k+1} \in \begin{cases} [\max(\sigma_{\min}, \gamma_1 \sigma_k), \sigma_k] & \text{if } \rho_k \geq \eta_2, \\ [\sigma_k, \gamma_2 \sigma_k] & \text{if } \rho_k \in [\eta_1, \eta_2), \\ [\gamma_2 \sigma_k, \gamma_3 \sigma_k] & \text{if } \rho_k < \eta_1. \end{cases} \quad (5.5)$$

Step 5: Relative accuracy update. Set

$$\omega_{k+1} = \min \left[\kappa_\omega, \frac{1}{\sigma_{k+1}} \right]. \quad (5.6)$$

Increment k by one and go to Step 1.

Theorem 4.6 then guarantees that the AR1DA Algorithm will find an ϵ -approximate first-order minimizer for the unconstrained version of problem (2.1) in at most $O(\epsilon^{-2})$ iterations (which is proved in [9] to be optimal) and at most $O(|\log(\epsilon)|\epsilon^{-2})$ approximate evaluations of the objective function and gradient. Note that

1. the accuracy requirement is truly adaptive and the absolute accuracy $\varepsilon_{1,i}$ may remain quite large as long as $\|\overline{\nabla_x^1 f}(x_k)\|$ itself remains large, as shown by item 3 in Step 1.
2. The accuracy requirement for computing \overline{f} does not depend on the absolute accuracy of the gradient, but only on its norm (squared). At initial iterations, this may be quite large.
3. In the unconstrained case, the AR1DA Algorithm is very close in spirit to the trust-region with dynamic accuracy of [13, Sections 8.4.1.1 and 10.6] and, when values of f are computed exactly, essentially recovers a proposal in [6]. It is also close to the proposal of [21], which is based on an Armijo-like linesearch and has similar accuracy requirements.

We now turn to the case where one seeks a first-order critical point for an unconstrained problem using approximate gradients and Hessians (under the assumption that the exact Hessian is Lipschitz continuous). As for the case $p = q = 1$, we have that (5.1) holds, making the verification of optimality in Step 1 relatively easy. Computing s_k is now more complicated but still practical, as it now implies minimizing the regularized quadratic model m_k starting from x_k until a step s_k is found such that

$$\|s_k\| \geq \mu\epsilon^{\frac{1}{2}} \quad \text{or} \quad \overline{\phi}_{m_k,1}^\delta(x_k + s_k, \omega_k) = \|\nabla_s^1 m_k(s_k)\| \leq \frac{1}{2}\theta\|s_k\|^2$$

(as proposed in [8], see also [17, 20, 7, 14]), with the additional constraint that, for $s_k \neq 0$,

$$\max[\varepsilon_{1,i}, \varepsilon_{2,i}] \leq \min \left[\frac{\omega_k \epsilon}{2}, \omega_k \frac{\overline{\Delta T}_2^f(x_k, s_k, \omega_k)}{\chi_2(\|s_k\|)} \right] \quad (5.7)$$

where

$$\overline{\Delta T}_2^f(x_k, s_k, \omega_k) = -\overline{\nabla_x^1 f}(x_k)^T s_k - \frac{1}{2} s_k^T \overline{\nabla_x^2 f}(x_k) s_k.$$

The resulting algorithm AR2DA is quite similar to AR1DA and is omitted for brevity. We note that

1. Algorithm AR2DA is guaranteed by Theorem 4.6 to find an ϵ -approximate first-order minimizer for the unconstrained version of problem (2.1) in at most $O(\epsilon^{-3/2})$ iterations (which is proved in [9] to be optimal) and at most $O(|\log(\epsilon)|\epsilon^{-3/2})$ approximate evaluations of the objective function, gradient and Hessian.
2. As for AR1DA, the absolute accuracies required by AR2DA on the approximate function, gradient and Hessian only depend on the magnitude of the Taylor increment, which is typically quite large in early iterations. The relative errors on the latter two remain bounded away from zero.
3. The absolute accuracies required on the approximate gradient and Hessian are comparable in magnitude, although (5.7) could be exploited to favour one with respect to the other.

The case where $p = 2$ and $q = 2$ (i.e. when second-order solutions are sought) is also computationally quite accessible: calculating the optimality measure $\bar{\phi}_{f,1}^{\delta_k}(x_k, \omega_k)$ or $\bar{\phi}_{m_k,1}^{\delta_k}(x_k + s_k, \omega_k)$ now involve a standard trust-region subproblem, for which both exact and approximate numerical solvers are known (see [13, Chapter 7] for instance), but the rest of the algorithm — in particular its adaptive accuracy requirement — is very similar to what we just discussed (see also [9]). Theorem 4.6 then ensures that resulting method converges to an ϵ -approximate second-order-necessary minimizer for the unconstrained version of problem (2.1) in at most $O(\epsilon^{-3})$ iterations and at most $O(|\log(\epsilon)|\epsilon^{-3})$ approximate evaluations of the objective function, gradient and Hessian.

We conclude this section by a brief discussion of the case where $q = 1$ and $p \in \{1, 2\}$ as before, but where \mathcal{F} is now defined by bound constraints. It is clear that evaluating and enforcing such constraints (by projection, say) has negligible cost and therefore falls in our framework. In this case, the calculations of $\bar{\phi}_{f,1}^{\delta_k}(x_k, \omega_k)$ or $\bar{\phi}_{m_k,1}^{\delta_k}(x_k + s_k, \omega_k)$ now involve simple linear optimization problems⁽¹⁾, which is computationally quite tractable. If $p = 1$, Step 2.2 and 2.3 involve convex quadratic optimization, while they involve minimizing a regularized quadratic model if $p = 2$. All results remain the same, and the ARpDA algorithm is then guaranteed to find a bound-constrained approximate first-order approximate minimizer in at most $O(\epsilon^{-2})$ or $O(\epsilon^{-3/2})$ iterations (which is proved in [9] to be optimal) and at most $O(|\log(\epsilon)|\epsilon^{-2})$ or $O(|\log(\epsilon)|\epsilon^{-3/2})$ approximate evaluations of the objective function, gradient and Hessian. The same algorithms and results obviously extend to the case where \mathcal{F} is a convex polyhedral set or any closed non-empty convex set, provided the cost of the projection on this set remains negligible compared to that of (approximately) evaluating the objective function and its derivatives.

6 A stochastic viewpoint on ARpDA

6.1 Probabilistic complexity

In this section we consider the case where the bounds $\{\varepsilon_j\}_{j=1}^p$ on the absolute errors on the derivative tensors $\{\nabla_x^j f(x)\}_{j=1}^p$ are satisfied with probability at least $(1 - t)$, with $t \in (0, 1)$. This may occur, for instance, if the approximate derivative tensors are obtained by some stochastic sampling scheme, as we detail below. We therefore assume that

$$Pr \left[\|\bar{\nabla}_x^j f(x_k) - \nabla_x^j f(x_k)\|_{[j]} \leq \varepsilon_j \right] \geq (1 - t) \quad \text{for each } j \in \{1, \dots, p\}. \quad (6.1)$$

We also assume that inequalities (2.21) and (2.22) in Step 3 of the ARpDA algorithm are satisfied with probability at least $(1 - t)$, i.e.

$$Pr \left[|\bar{f}_k(x_k + s_k, \omega_k) - f(x_k + s_k)| \leq \varepsilon_0 \right] \geq 1 - t, \quad (6.2)$$

and

$$Pr \left[|\bar{f}_k(x_k, \omega_k) - f(x_k)| \leq \varepsilon_0 \right] \geq 1 - t \quad (6.3)$$

where we have defined $\varepsilon_0 \stackrel{\text{def}}{=} \omega_k |\bar{\Delta T}_p^f(x_k, s_k, \omega_k)|$. Clearly, different values for t could be chosen in (6.1), one for each index (tensor order) $j \in \{1, \dots, p\}$. Similarly, different values of

⁽¹⁾Formerly known as linear programming problems, or LPs.

t in (6.2) and (6.3) could be considered. However, for the sake of simplicity, we assume here that all the inequalities involved in (6.1)–(6.3) hold with the same fixed lower bound $(1 - t)$ on the probability of success. We also assume that the events in (6.1)–(6.3) are independent.

We stress that Algorithms 3.2 and 3.3 terminates independently of the satisfaction of the accuracy requirements on the tensor derivatives. This is due to the fact that termination relies on the inequality (3.5). Moreover, during the iterations of either of these algorithms before the last, it may happen that the accuracy on the tensor derivatives fails to be achieved, but this has not impact on the worst-case complexity. Satisfying the accuracy requirement is only crucial in the last iteration of Algorithm 3.2 or 3.3 (that is in Steps 1.2 and 2.2). Let $\mathcal{E}_r(S)$ be the event: “the relations

$$\|\overline{\nabla_x^j f(x_k)} - \nabla_x^j f(x_k)\|_{[j]} \leq \varepsilon_j \text{ for all } j \in \{1, \dots, r\}$$

hold for some j at Step S of the last iteration of the relevant algorithm”. In Step 1.2, inexact values are computed for the first q derivatives, and the probability that event $\mathcal{E}_q(1.2)$ occurs is therefore at least $(1 - t)^q$. Similarly, the probability that event $\mathcal{E}_q(2.2)$ occurs is at least $(1 - t)^p$. Finally, at Step 3 of the ARpDA algorithm, the probability that both (2.21) and (2.22) hold is at least $(1 - t)^2$. Then, letting for $i \in \{1, \dots, k\}$, $\mathcal{E}_{[i]}$ be the event: “Inequalities (2.13), (2.21) and (2.22) hold at iteration i , of the ARpDA algorithm”, the probability that $\mathcal{E}_{[i]}$ occurs is then at least $(1 - t)^{p+q+2}$. Finally, letting $\mathcal{E}(k)$ be the event: “ $\mathcal{E}_{[i]}$ occurs for all iterations $i \in \{1, \dots, k\}$ of the ARpDA algorithm”, we deduce that

$$Pr[\mathcal{E}(k)] \equiv Pr\left[\bigcap_{i=1}^k \mathcal{E}_{[i]}\right] \geq (1 - t)^{k(p+q+2)}.$$

Thus, requiring that the event $\mathcal{E}(k)$ occurs with probability at least $1 - \bar{t}$, we obtain that

$$Pr[\mathcal{E}(k)] \geq (1 - t)^{k(p+q+2)} = 1 - \bar{t}, \text{ i.e., } t = 1 - (1 - \bar{t})^{\frac{1}{k(p+q+2)}} = O\left(\frac{\bar{t}}{k(p+q+2)}\right).$$

Taking into account that, when (2.13), (2.21) and (2.22) hold, the ARpDA algorithm terminates in at most $k = O(\epsilon^{-\frac{p+\beta}{p-q+\beta}})$ iterations (as stated by Theorem 4.6), we deduce the following result.

Theorem 6.1 Let $f \in \mathcal{C}^{p,\beta}(\mathbb{R}^n)$. Suppose that the probabilistic assumptions of this section hold and that, at each of iteration of the ARpDA algorithm, the probability t satisfies

$$t = O\left(\frac{\bar{t} \epsilon^{\frac{p+\beta}{p-q+\beta}}}{(p+q+2)}\right). \quad (6.4)$$

Then, given $\epsilon \in (0, 1)$, the conclusions of Theorem 4.6 hold with probability at least $(1 - \bar{t})$.

As a consequence, when $p = q = 2$ and $\beta = 1$ we have to choose $t = O\left(\frac{\bar{t}\epsilon^3}{6}\right)$, while, when $p = q = \beta = 1$, we have to choose $t = O\left(\frac{\bar{t}\epsilon^2}{3}\right)$.

We stress that the above analysis is unduly pessimistic in the case where $p = q = 1$. Indeed, as already noticed in Section 5, no reduction in $\{\varepsilon_j\}$ is necessary at Step 2, as (2.13) is automatically enforced whenever the relative error on the first derivative $\overline{\nabla_x^1 f}(x)$ is bounded by ω . Noting that this last event has probability at least $1 - t$, we can conclude that $Pr(\mathcal{E}_{[i]}) \geq (1 - t)^3$ and to get the optimal complexity $O(\epsilon^{-2})$ with probability at least $1 - \bar{t}$, we need to choose $t = O\left(\frac{\bar{t}\epsilon^2}{3}\right)$. We also emphasize that the purpose of Theorem 6.1 is limited to offer guidance on desirable value of t and not to prescribe an algorithmically binding bound. Indeed some of the constants involved in the bound of Theorem 4.6 (and thus of Theorem 6.1) are typically unknown a priori (which is why we have not been more specific in (6.4)).

6.2 Sample size in subsampling for finite-sum problems

In what follows, we now focus on the solution of large-scale instances of the finite-sum problems arising in machine learning and data analysis, that are modelled as

$$\min_{x \in \mathbb{R}^n} f(x) = \frac{1}{N} \sum_{i=1}^N \psi_i(x), \quad (6.5)$$

with $N > 0$ and $\psi_i : \mathbb{R}^n \rightarrow \mathbb{R}$. Restricting ourselves to the cases where $p \leq 2$, we discuss the application of Algorithm AR1DA and AR2DA to problem (6.5). In this case, the approximation of the objective function's value and of first and second derivatives is obtained by a subsampling procedures, i.e. these quantities are approximated by randomly sampling component functions ψ_i . More precisely, at iteration k these approximations take the form:

$$\begin{aligned} \bar{f}_k(x_k) &= \frac{1}{|\mathcal{D}_k|} \sum_{i \in \mathcal{D}_k} \psi_i(x_k), & \bar{f}_k(x_k + s_k) &= \frac{1}{|\mathcal{D}_k|} \sum_{i \in \mathcal{D}_k} \psi_i(x_k + s_k), \\ \overline{\nabla_x^1 f}(x_k) &= \frac{1}{|\mathcal{G}_k|} \sum_{i \in \mathcal{G}_k} \overline{\nabla_x^1 \psi_i}(x_k), & \text{and} & \quad \overline{\nabla_x^2 f}(x_k) = \frac{1}{|\mathcal{H}_k|} \sum_{i \in \mathcal{H}_k} \overline{\nabla_x^2 \psi_i}(x_k), \end{aligned}$$

where \mathcal{D}_k , \mathcal{G}_k and \mathcal{H}_k are subsets of $\{1, 2, \dots, N\}$. The question then arises of estimating the cardinality of these sample sets in order to ensure that the approximations of the objective function's value and its first and second derivatives satisfy (6.1) for $j = 1$ and $j = 2$, (6.2) and (6.3). This is the object of our final result, whose proof hinges on the operator-Bernstein inequality.

Theorem 6.2 Suppose that there exist non-negative constants $\{\kappa_{\psi,j}\}_{j=0}^2$ such that, for $x \in \mathbb{R}^n$ and all $j \in \{0, 1, 2\}$

$$\max_{i \in \{1, \dots, N\}} \|\nabla_x^j \psi_i(x)\| \leq \kappa_{\psi,j}(x). \quad (6.6)$$

Suppose that a subsample \mathcal{A}_k is chosen randomly and uniformly from $\{1, \dots, N\}$ and that, for some $j \in \{0, 1, 2\}$, one computes

$$\overline{\nabla_x^j f}(x) = \frac{1}{|\mathcal{A}_k|} \sum_{i \in \mathcal{A}_k} \overline{\nabla_x^j \psi_i}(x),$$

with

$$|\mathcal{A}_k| \geq \frac{4\kappa_{\psi,j}(x)}{\varepsilon_j} \left(\frac{2\kappa_{\psi,j}(x)}{\varepsilon_j} + \frac{1}{3} \right) \log \left(\frac{d}{t} \right), \quad (6.7)$$

where

$$d = \begin{cases} 2, & \text{if } j = 0, \\ n + 1, & \text{if } j = 1, \\ 2n, & \text{if } j = 2. \end{cases}$$

Then, given $t \in (0, 1)$, condition (6.1) holds for $x = x_k$ with probability at least $(1 - t)$ if $j \in \{1, 2\}$, or, if $j = 0$, each of the conditions (6.2) and (6.3) holds with probability at least $(1 - t)$ for $x = x_k + s_k$ and $x = x_k$, respectively.

Proof. Let $\mathcal{A}_k \subseteq \{1, \dots, N\}$ be a sample set of cardinality $|\mathcal{A}_k|$. Consider $j \in \{0, 1, 2\}$ and $|\mathcal{A}_k|$ random tensors $Z_u(x)$ such that,

$$\Pr [Z_u(x) = \nabla_x^j \psi_i(x)] = \frac{1}{N}, \quad (i \in \{1, \dots, N\}).$$

For $u \in \mathcal{A}_k$, let us define

$$X_u \stackrel{\text{def}}{=} (Z_u(x) - \nabla_x^j f(x)), \quad \overline{\nabla_x^j f(x)} \stackrel{\text{def}}{=} \frac{1}{|\mathcal{A}_k|} \sum_{u \in \mathcal{A}_k} Z_u(x)$$

and

$$X \stackrel{\text{def}}{=} \sum_{u \in \mathcal{A}_k} X_u = |\mathcal{A}_k| \left(\overline{\nabla_x^j f(x)} - \nabla_x^j f(x) \right).$$

Since (6.5) gives that

$$\frac{1}{N} \sum_{i=1}^N \nabla_x^j \psi_i(x) = \nabla_x^j f(x),$$

we deduce that

$$E(X_u) = \frac{1}{N} \sum_{i=1}^N (\nabla_x^j \psi_i(x) - \nabla_x^j f(x)) = 0, \quad u \in \mathcal{A}_k.$$

Moreover, assuming $Z_u(x) = \nabla_x^j \psi_l(x)$ for some $l \in \{1, \dots, N\}$ and using (6.6), we have that

$$\|X_u\| \leq \left\| \frac{N-1}{N} \nabla_x^j \psi_l(x) - \frac{1}{N} \sum_{i \in \{1, \dots, N\} \setminus \{l\}} \nabla_x^j \psi_i(x) \right\| \leq 2 \frac{N-1}{N} \kappa_{\psi,j}(x) \leq 2\kappa_{\psi,j}(x),$$

so that the variance of X can be bounded as follows:

$$\begin{aligned}
v(X) &= \max [\|E(XX^T)\|, \|E(X^T X)\|] \\
&= \max \left[\left\| \sum_{u \in \mathcal{A}_k} E(X_u X_u^T) \right\|, \left\| \sum_{u \in \mathcal{A}_k} E(X_u^T X_u) \right\| \right] \\
&\leq \max \left[\sum_{u \in \mathcal{A}_k} \|E(X_u X_u^T)\|, \sum_{u \in \mathcal{A}_k} \|E(X_u^T X_u)\| \right] \\
&\leq \max \left[\sum_{u \in \mathcal{A}_k} E(\|X_u X_u^T\|), \sum_{u \in \mathcal{A}_k} E(\|X_u^T X_u\|) \right] \\
&\leq \sum_{u \in \mathcal{A}_k} E(\|X_u\|^2) \leq 4|\mathcal{A}_k| \kappa_{\psi,j}^2(x),
\end{aligned}$$

in which the first and the third inequalities hold because of the triangular inequality, while the second is due to the Jensen's inequality (note that the spectral norm $\|\cdot\|$ is convex). Therefore, according to the Operator-Bernstein Inequality stated in [23, Theorem 6.1.1], we obtain that

$$Pr \left[\|\overline{\nabla_x^j f(x)} - \nabla_x^j f(x)\| \geq \epsilon_j \right] = Pr \left[\|X\| \geq \epsilon_j |\mathcal{A}_k| \right] \leq d e^{-\frac{\epsilon_j^2 |\mathcal{A}_k|}{4\kappa_{\psi,j}(x) \left(2\kappa_{\psi,j}(x) + \frac{1}{3}\epsilon_j \right)}}, \quad (6.8)$$

with $d = 2$ if $j = 0$, $d = n + 1$ if $j = 1$ and $d = 2n$ if $j = 2$. Then, bounding the right-hand side of (6.8) by t , taking logarithms and extracting $|\mathcal{A}_k|$ gives (6.7). \square

In particular, Theorem 6.2 gives the lower bounds

$$|\mathcal{D}_k| \geq \frac{4\kappa_{\psi,j}(x)}{\varepsilon_0} \left(\frac{2\kappa_{\psi,j}(x)}{\varepsilon_0} + \frac{1}{3} \right) \log \left(\frac{2}{t} \right), \quad |\mathcal{G}_k| \geq \frac{4\kappa_{\psi,j}(x)}{\varepsilon_1} \left(\frac{2\kappa_{\psi,j}(x)}{\varepsilon_1} + \frac{1}{3} \right) \log \left(\frac{n+1}{t} \right)$$

and

$$|\mathcal{H}_k| \geq \frac{4\kappa_{\psi,j}(x)}{\varepsilon_2} \left(\frac{2\kappa_{\psi,j}(x)}{\varepsilon_2} + \frac{1}{3} \right) \log \left(\frac{2n}{t} \right).$$

The adaptive nature of these sample sizes is apparent in these formulae, because they depend on x and ε_j , which are themselves dynamically updated in the course of the ARpDA algorithm. If only global information is available on the size of the j -th derivative of the ψ_i , the dependence on x may obviously be avoided by choosing a uniform upper bound $\kappa_{\psi,j}$ for all $x \in \mathcal{F}$, at the cost of a lesser adaptivity.

We finally emphasize that the per-iteration failure probability t given in (6.4) is not too demanding in what concerns the sample size, because it only occurs in the logarithm term of (6.7). The same is true of the impact of the value of the unknown constants hidden in the $O(\cdot)$ notation in (6.4).

7 Conclusion and perspectives

We have provided a general regularization algorithm using inexact function and derivatives' values, featuring a flexible adaptive mechanism for specifying the amount of inexactness acceptable at each iteration. This algorithm, inspired by the unifying framework proposed in [9], is applicable to unconstrained and inexpensively-constrained nonconvex optimization problems, and provides optimal iteration complexity for arbitrary degree of available derivatives,

arbitrary order of optimality and the full range of smoothness assumptions on the objective function highest derivative. We have also specialized this algorithm to the cases of first- and second-order methods, exhibiting simple and numerically realistic methods. We have finally provided a probabilistic version of the complexity analysis and derived associated lower bounds on sample size in the context of subsampling methods.

There are of course many ways in which the proposed algorithm might be improved. For instance, the central calculation of relatively accurate Taylor increments may possibly be made more efficient by updating the absolute accuracies for different degrees separately. Further techniques to avoid unnecessary derivative computations (without affecting the optimal complexity) could also be investigated.

The framework proposed in this paper also offers obvious avenues for specializations to specific contexts, among which we outline two. The first is that of algorithms using stochastic approximations of function values and derivatives. The technique presented here derives probabilistic conditions under which properties of the deterministic algorithms are preserved. It does not provide an algorithm which is robust against failures to satisfy the adaptive accuracy requirements. This is in contrast with the interesting analysis of unconstrained first-order methods of [21] and [5]. Combining the generality of our approach with the robustness of the proposal in these latter papers is thus desirable. The second interesting avenue is the application of the new results to multi-precision optimization in the context of very high performance computing. In this context, it is of paramount importance to limit energy dissipation in the course of an accurate calculation, and this may be obtained by varying the accuracy of the most crucially expensive of its parts (see [16] for unconstrained quadratic optimization). The discussion above again provides guidance at what level of arithmetic accuracy is needed to achieve overall performance while maintaining optimal complexity. Both these topics are the object of ongoing research and will be reported on at a later stage.

Acknowledgment

INdAM-GNCS partially supported the first and third authors under Progetti di Ricerca 2018. The last author gratefully acknowledges the support and friendly environment provided by the Department of Industrial Engineering at the Università degli Studi, Florence (Italy) during his visit in the fall of 2018.

References

- [1] S. Bellavia, S. Gratton, and E. Riccietti. A Levenberg-Marquardt method for large nonlinear least-squares problems with dynamic accuracy in functions and gradients. *Numerische Mathematik*, 140:791–825, 2018.
- [2] S. Bellavia, G. Gurioli, and B. Morini. Theoretical study of an adaptive cubic regularization method with dynamic inexact Hessian information. arXiv:1808.06239, 2018.
- [3] E. Bergou, Y. Diouane, V. Kungurtsev, and C. W. Royer. A subsampling line-search method with second-order results. arXiv:1810.07211, 2018.
- [4] E. G. Birgin, J. L. Gardenghi, J. M. Martínez, S. A. Santos, and Ph. L. Toint. Worst-case evaluation complexity for unconstrained nonlinear optimization using high-order regularized models. *Mathematical Programming, Series A*, 163(1):359–368, 2017.
- [5] J. Blanchet, C. Cartis, M. Menickelly, and K. Scheinberg. Convergence rate analysis of a stochastic trust region method via supermartingales. arXiv:1609.07428v3, 2018.
- [6] T. Bonniot. Convergence and complexity of unconstrained optimization methods with inexact gradients. Master’s thesis, ENSEEIHT, Toulouse, France, September 2018. (supervised by S. Gratton, D. Orban and Ph. Toint).
- [7] C. Cartis, N. I. M. Gould, and Ph. L. Toint. Trust-region and other regularization of linear least-squares problems. *BIT*, 49(1):21–53, 2009.

- [8] C. Cartis, N. I. M. Gould, and Ph. L. Toint. Adaptive cubic overestimation methods for unconstrained optimization. Part II: worst-case function-evaluation complexity. *Mathematical Programming, Series A*, 130(2):295–319, 2011.
- [9] C. Cartis, N. I. M. Gould, and Ph. L. Toint. Sharp worst-case evaluation complexity bounds for arbitrary-order nonconvex optimization with inexpensive constraints. arXiv:1811.01220, 2018.
- [10] C. Cartis, N. I. M. Gould, and Ph. L. Toint. Worst-case evaluation complexity and optimality of second-order methods for nonconvex smooth optimization. To appear in the Proceedings of the 2018 International Conference of Mathematicians (ICM 2018), Rio de Janeiro, 2018.
- [11] C. Cartis and K. Scheinberg. Global convergence rate analysis of unconstrained optimization methods based on probabilistic models. *Mathematical Programming, Series A*, 159(2):337–375, 2018.
- [12] X. Chen, B. Jiang, T. Lin, and S. Zhang. On adaptive cubic regularization Newton’s methods for convex optimization via random sampling. arXiv:1802.05426, 2018.
- [13] A. R. Conn, N. I. M. Gould, and Ph. L. Toint. *Trust-Region Methods*. MPS-SIAM Series on Optimization. SIAM, Philadelphia, USA, 2000.
- [14] J. P. Dussault. Simple unified convergence proofs for the trust-region and a new ARC variant. Technical report, University of Sherbrooke, Sherbrooke, Canada, 2015.
- [15] S. Gratton, A. Sartenaer, and Ph. L. Toint. Recursive trust-region methods for multiscale nonlinear optimization. *SIAM Journal on Optimization*, 19(1):414–444, 2008.
- [16] S. Gratton, E. Simon, and Ph. L. Toint. Minimizing convex quadratics with variable precision Krylov methods. arXiv:1807.07476, 2018.
- [17] A. Griewank. The modification of Newton’s method for unconstrained optimization by bounding cubic terms. Technical Report NA/12, Department of Applied Mathematics and Theoretical Physics, University of Cambridge, Cambridge, United Kingdom, 1981.
- [18] L. Liu, X. Liu, C.-J. Hsieh, and D. Tao. Stochastic second-order methods for non-convex optimization with inexact Hessian and gradient. arXiv:1809.09853, 2018.
- [19] Yu. Nesterov. *Introductory Lectures on Convex Optimization*. Applied Optimization. Kluwer Academic Publishers, Dordrecht, The Netherlands, 2004.
- [20] Yu. Nesterov and B. T. Polyak. Cubic regularization of Newton method and its global performance. *Mathematical Programming, Series A*, 108(1):177–205, 2006.
- [21] C. Paquette and K. Scheinberg. A stochastic line search method with convergence rate analysis. arXiv:1807.07994, 2018.
- [22] N. Tripuraneni, M. Stern, J. Regier, and M. I. Jordan. Stochastic cubic regularization for fast nonconvex optimization. arXiv:1711.02838v2, 2017.
- [23] J. Tropp. *An Introduction to Matrix Concentration Inequalities*. Number 8:1-2 in Foundations and Trends in Machine Learning. Now Publishing, Boston, USA, 2015.
- [24] S. A. Vavasis. Black-box complexity of local minimization. *SIAM Journal on Optimization*, 3(1):60–80, 1993.
- [25] P. Xu, F. Roosta-Khorasani, and M. W. Mahoney. Newton-type methods for non-convex optimization under inexact Hessian information. arXiv:1708.07164v3, 2017.
- [26] P. Xu, F. Roosta-Khorasani, and M. W. Mahoney. Second-order optimization for non-convex machine learning: An empirical study. arXiv:1708.07827v2, 2018.
- [27] Z. Yao, P. Xu, F. Roosta-Khorasani, and M. W. Mahoney. Inexact non-convex Newton-type methods. arXiv:1802.06925, 2018.